



# What Archie is, and what it does on the computer it runs on

---

A plain description of the software for anyone who has been asked to approve it: where it installs, what it sends over the network, where credentials are kept, and what it is not allowed to do.

---

#### PRODUCT

**Archie, version 0.2.3 (beta)**

#### REVISION

**5, September 3, 2026**

#### PUBLISHER

**Otian AI LLC**

#### CONTACT

**jett@otianai.com  
jack@otianai.com**

#### PLATFORMS

**macOS and Windows**

#### WEBSITE

**otianai.com**

## SECTION 1

# What Archie is

---

Archie is a desktop application that lets one person set up an AI assistant and then talk to it from a chat app they already use.

It is installed like any other desktop program and it runs only while it is open. There is no Archie server that holds the assistant: the assistant runs on the computer where the app is installed. Close the app and it stops answering.

The person using it does three things. First they connect an **AI account**. This is an API key from a provider such as Anthropic, OpenAI, or Google: a developer account billed by what it is used for, not a consumer chat subscription like ChatGPT Plus or Claude Pro. The difference matters for a review and section 16 sets it out, but the short form is that the major providers do not train on what is sent through their API, and on an ordinary paid key that is the default rather than a setting somebody has to find. There is one alternative to connecting a key, and a reviewer should know it exists before reading further: Archie is also sold on a plan that includes the AI, and on that plan the calls run on the vendor's own account and pass through the vendor's server. Section 8 is the whole of what that changes. Second, they optionally connect a chat app: Telegram, Slack, Discord, Matrix, or, on a Mac, the Messages app, so they can reach the assistant from their phone in addition to their computer. Third, they install **skills**, which are small bundles of instructions that tell the assistant how to handle one kind of job, like drafting a reply to an email or putting an event on a calendar.

What the assistant can actually touch is decided entirely by what the user connects. A fresh install can do nothing except talk. Reading a mailbox requires that person to sign in to that mailbox and grant permission. Reading a calendar is a separate permission. Nothing is on by default and nothing is inherited from another user.

### THE SHORT VERSION FOR A REVIEWER

Archie is a locally installed client. It makes outbound HTTPS connections to the services the user connected and to the vendor's own backend. It opens no listening network port, accepts no inbound connection, and installs no service, driver, or kernel extension. It runs no shell; the few programs it does start, such as a browser already on the computer or a media tool it downloaded and checked against a published hash, are named in section 10. It does not require administrator rights to install on Windows.

It is also an AI agent that reads text other people wrote, and that has a class of weakness no vendor has solved. Section 11 is about that, and it is the section to read first if you only read one.

## SECTION 2

# What this document claims, and what it does not

---

A security claim means nothing without the boundary it holds inside. This section is that boundary, stated before anything else, so every sentence after it can be checked against something.

## What Archie is built to defend against

- **A remote attacker on the network.** There is nothing to connect to. No listening port, no inbound path, no service.
- **Model output being turned into execution.** No shell, no eval, no dynamic code loading, anywhere.
- **An assistant being talked into reaching your internal network.** On every path that goes through Archie's own HTTP client, private, loopback, and carrier-grade NAT addresses are refused, and the refusal happens inside the name resolver. Section 10 names the two lanes that sit outside it and what bounds each.
- **Secrets sitting in readable places.** Credentials live in the operating system credential store and are read transiently. They are never in the database, in logs, or in exports.
- **A malicious or broken add-on introducing a binary.** An add-on can reference a vetted, hash-pinned binary by identifier. It can never name a URL.
- **An action being taken without a person.** Sending mail, changing a calendar, and writing to a connected service all stage as a proposal and commit on a human action.

## What it does not defend against, stated plainly

| NOT DEFENDED                                   | WHAT THAT MEANS FOR YOU                                                                                                                                                                                                           |
|------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| An endpoint that is already compromised        | Archie assumes the computer it runs on is trusted. Anything already running as that user can read Archie's database, its files, and, on Windows, its credentials. Nothing in this document survives a compromised machine.        |
| A local attacker inside the user's own session | The audit log, the credential store, and the update path are all reachable by anything running as that user. See sections 10, 15, and 17.                                                                                         |
| The user themselves                            | Archie is the user's tool. It does not stop them from connecting an account, approving a draft, or pasting a key. There is no administrator above them in the personal edition, and the license gates are on a computer they own. |
| Prompt injection, fully                        | The gate stops an injected instruction from changing things silently. It does not stop it from leaking. This is the real one, and section 11 is the whole story rather than a line here.                                          |
| What your AI provider does                     | Content goes to the provider whose key was pasted, under your agreement with them. Archie is not party to it. Section 16.                                                                                                         |
| What a chat platform does                      | If a chat app is connected, the assistant's messages transit that platform on its terms. Section 7 covers what that actually exposes.                                                                                             |

### ONE RULE THIS DOCUMENT TRIES TO FOLLOW

Where it says "everything" or "the whole list", it means it, and you should check it. Where it cannot be complete it says so instead. If you find something Archie does that is not described here, that is a defect in this document and worth telling us about, because the value of a page like this is entirely in whether it survives being checked.

## SECTION 3

# How the application is built

---

Archie is three parts, each allowed to do less than the one below it. The split exists so that the part handling untrusted text is the part furthest from the credentials.

### The window

NO SECRETS

The screens people look at, written in React and TypeScript and drawn by the operating system's own web view. It holds no credentials, runs no database queries, and cannot read or write files. It asks the layer below for everything, through a fixed list of named commands that is deny-by-default.

↓ **typed commands**, allowlisted one by one

### The core

HOLDS THE KEYS

Rust. Owns the local database, the operating system credential store, and the audit log. This is the only part that can read a secret value, and it reads one only at the moment it is putting it into a specific outgoing request. A secret is never handed back up to the window after it is typed in.

↓ **direct calls inside one process**, no network port

### The runtime

RUNS THE ASSISTANT

The part that does the actual work: deciding which skill a message belongs to, calling the AI provider, and running tools. Each assistant runs as its own supervised task, so one that hangs or fails is restarted on its own without taking the app down with it. It makes the outbound calls a running assistant needs, and it accepts no inbound connection.

The application is written in Rust and TypeScript and packaged with Tauri 2, which means it uses the operating system's built-in web view rather than shipping a copy of a browser.

### WHAT THAT PICTURE IS, AND WHAT IT IS NOT

The three layers above are a **code** boundary, enforced by which crate may call what and by the allowlist between the window and the core. They are not three operating-system processes. Archie ships as one process: the window is the system web view, and the core and the runtime are Rust inside the same binary, with each running assistant supervised as its own task so it can be restarted alone.

The consequence, said plainly rather than left for a reviewer to find: the code that handles untrusted text runs in the same process that can read the credential store. What keeps a secret out of it is the rule that a credential is resolved at the moment it is put into one specific outgoing request, not a memory boundary the operating system enforces. An earlier version of this document described the runtime as a separate child process talking over a pipe. That was the plan and it is not what ships.

**One note about that word deny-by-default.** It is accurate for the boundary above, where the window may only call commands that were explicitly granted. Until August 2026 it was not true of the trial mail gate, which was a denylist on tool-name prefixes while this document described it as deny-by-default. That gate was rewritten rather than the sentence softened; section 8 has the detail, and it is named here because a reviewer is entitled to know which claims have been wrong before.

## What gets installed, and where

| PLATFORM | INSTALLER                                                                            | RIGHTS NEEDED                                                                                                                      |
|----------|--------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------|
| macOS    | A signed disk image ( <code>.dmg</code> ). The user drags the app into Applications. | None beyond the ability to write to Applications. It can also be run from the user's own folder.                                   |
| Windows  | An NSIS installer ( <code>.exe</code> ), built in <b>current user</b> mode.          | None. It installs under the signed-in user's own profile and does not write to Program Files or to the machine-wide registry hive. |

No background service, scheduled task, daemon, driver, or kernel extension is installed. Nothing runs at startup unless the user turns that on. When the app is closed, no Archie process remains.

**The Windows install location is a tradeoff, not a free win.** Needing no administrator rights is convenient and it also means the installed executable sits somewhere the user's own account can write. Section 17 says what follows from that.

### Files it creates

| WHAT                                                 | WHERE ON MACOS                                                                                                                                                           | HOW IT IS STORED                                                                                                                                             |
|------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Settings, agent configuration, conversation history  | <code>~/Library/Application Support/com.archie.app/archie.db</code>                                                                                                      | A SQLite file. Not separately encrypted. See the note below.                                                                                                 |
| Credentials and API keys                             | macOS Keychain, service <code>com.archie.app</code>                                                                                                                      | Operating system credential store. Never written to the database or to any file.                                                                             |
| Logs                                                 | One per assistant, at <code>.../com.archie.app/workspaces/&lt;workspace&gt;/agents/&lt;assistant&gt;/logs/</code> , plus a crash log in <code>.../com.archie.app/</code> | JSON lines, rotated at about 512 KB so they cannot grow without limit. Every line passes a redactor that strips secret-shaped values before it reaches disk. |
| Files the assistant saved or downloaded for the user | <code>~/Downloads/Archie/</code>                                                                                                                                         | Ordinary files, in the user's own Downloads folder where they can see them.                                                                                  |
| Helper programs and models an add-on needed          | <code>~/Library/Application Support/com.archie.app/assets/</code>                                                                                                        | Downloaded from the vendor's asset host and checked against a published SHA-256 before use (section 10).                                                     |

| WHAT                                                                                                                                                                                                  | WHERE ON MACOS                                                        | HOW IT IS STORED                                                                                             |
|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------|
| Each assistant's own folder: its instructions, its skills and their settings, its memory, any document the user gave it, and, if a website job was set up, the browser profile that job signs in with | <code>~/Library/Application Support/com.archie.app/workspaces/</code> | Ordinary files. Not separately encrypted.                                                                    |
| A weekly backup, only if the user switches it on                                                                                                                                                      | A folder the user picks, anywhere they can write                      | A plain zip, not encrypted, holding a copy of the database and every assistant's folder. See the note below. |

On Windows the same folders live under the user's `AppData` and Downloads folders, and credentials live in Windows Credential Manager instead of the Keychain. Those two stores are not equivalent, and section 17 says how.

#### THE BACKUP IS THE ONE FILE THAT CAN LEAVE THE FOLDERS ABOVE, AND IT IS WORTH A PARAGRAPH

Archie can save a backup once a week on its own. It is **off until somebody turns it on and picks a folder**, it writes `Archie-backup-<date>.archie` there, and it keeps the three most recent. The file is an ordinary zip with no encryption of its own, and what is inside it is the conversation history, every assistant's instructions, memory and documents, and, where a website job has been set up, that assistant's browser profile, which can hold a live session for a site somebody signed in to. No credential from the Keychain is in it; the file lists which keys an assistant used by label and last four characters, and a restore asks for them again.

Two consequences follow, and the second is the one to weigh. It never travels by itself, because nothing in that path opens a network connection. But the app suggests picking a folder that syncs to cloud storage, so that a backup survives the computer, and a person who takes that advice has put an unencrypted copy of the conversation history into whatever cloud account that folder belongs to. If your policy has a view about that, the controls are the switch, the choice of folder, and whole-disk encryption underneath it.

#### WORTH KNOWING

The local database is a plain SQLite file. It holds settings, the assistant's configuration, and conversation history, and anyone who can read that user's profile folder can read it. It never holds an API key, a password, or an OAuth token. If conversation history on the machine is a concern, the control is whole-disk encryption: FileVault on macOS or BitLocker on Windows.

## SECTION 5

# Network behavior

Every connection Archie makes is outbound HTTPS or a secure websocket that it opened itself. Nothing on the network can connect to it.

Inbound: none

Archie listens on no port and accepts no inbound connection. There is one narrow exception, and it lasts seconds:

- When the user signs in to Google, Microsoft, or their Otian account, the app opens a listener on `127.0.0.1` on a port the operating system picks, because that is how the standard OAuth desktop flow returns the user to the app after they approve in their browser. It is bound to the loopback address, so it is not reachable from the network, and it closes as soon as the browser hands back the response.
- **That flow uses PKCE** (`archie-core::oauth` generates a verifier and challenge per attempt) together with a state nonce that is validated on return. Both matter here: on a shared machine, a loopback redirect without PKCE is interceptable by any other local process, and PKCE is what makes an intercepted authorization code useless.

The part that executes the assistant opens the outbound connections it needs, to the AI provider and to any chat platform or service the user connected, and it listens on nothing. Every one of those destinations is in section 6.

Local network: only if the user asks for it

Archie has an optional feature for switching smart lights and plugs on the user's own wifi. It is off until the user runs a scan and adds a device by hand. If it is used, the app sends UDP discovery broadcasts on the local subnet, and it can then reach only the addresses that lane is allowed to reach, which are private ones. On a corporate network this feature has nothing to find and can simply be left alone; there is no need to use it.

One other thing can reach an address on the same network, and only because the user typed it: the AI provider setting accepts an address of the user's own, which is how somebody points Archie at a model running on a machine they operate (section 16). The model never names that address; the person does, once, in a settings field.

SECTION 6

Where it connects, and why

This is the full list of destinations, grouped by what causes the connection. Anything in the second group happens only because the user connected that service.

Always, as part of the application working

| HOST                                  | PURPOSE                                 | NOTES                                                                                                                                                                                                                                                                                                                                                         |
|---------------------------------------|-----------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>otianai.com</code>              | Sign-in, license check, update manifest | The vendor's own site. Also serves the optional phone page.                                                                                                                                                                                                                                                                                                   |
| <code>archie-4f35.onrender.com</code> | The vendor's backend service            | Sign-in, licensing, account management, account email, and relaying AI requests to Anthropic for the two configurations in section 8: the free starter credits, and the plan sold with the AI included. This is the host to add to a block list alongside <code>otianai.com</code> if you want to test that an agent on its own key keeps working without us. |

| HOST                                                              | PURPOSE                     | NOTES                                                                                                                                                                                                                                                                                    |
|-------------------------------------------------------------------|-----------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>*.googleapis.com</code>                                     | Firebase Auth and Firestore | The vendor's Firebase project holds the account record, the add-on catalog, the remote configuration document, and the license tier. Installing an add-on also increments a <b>global</b> counter on that add-on's catalog document; no per-person record of who installed what is kept. |
| <code>github.com</code>                                           | Downloading an app update   | Update binaries are GitHub release assets. Every update package is signature-verified before it is applied. Section 15 says what that does and does not prove.                                                                                                                           |
| <code>api.stripe.com</code> ,<br><code>checkout.stripe.com</code> | Subscription and billing    | Opened in the browser. Archie never sees a card number.                                                                                                                                                                                                                                  |

## Only after the user connects that service

| HOST                                                                                                                                                                                | WHAT CONNECTS TO IT                                                                                                                                                                                                                                                                                                                               |
|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>api.anthropic.com</code> , <code>api.openai.com</code> ,<br><code>generativelanguage.googleapis.com</code> , <code>api.groq.com</code> , and similar                          | The AI provider whose key the user pasted in. This is where the assistant's thinking happens.                                                                                                                                                                                                                                                     |
| <code>gmail.googleapis.com</code> , <code>www.googleapis.com</code> ,<br><code>tasks.googleapis.com</code> , <code>oauth2.googleapis.com</code>                                     | Gmail, Google Calendar, and Google Tasks, if the user signs in to a Google account.                                                                                                                                                                                                                                                               |
| <code>graph.microsoft.com</code> , <code>login.microsoftonline.com</code>                                                                                                           | Outlook mail and calendar, if the user signs in to a Microsoft 365 account.                                                                                                                                                                                                                                                                       |
| <code>api.telegram.org</code> , <code>discord.com</code> , <code>slack.com</code> , the user's Matrix homeserver                                                                    | Whichever chat app the user connected. Read section 7 before approving this one. The fifth channel, iMessage on a Mac, opens no connection of Archie's own: it reads and writes through the Messages app already signed in on that computer, so the traffic is Apple's and the account is the user's.                                             |
| <code>imap.mail.me.com</code> , <code>imap.fastmail.com</code> , and the equivalent for Yahoo, AOL, Zoho and GMX, with their <code>smtp</code> and <code>caldav</code> counterparts | Mail and calendar on one of those six providers, if the user connected it with an app password. The host table is fixed in the source: the user types an address and a password, never a server name. IMAP is TLS on 993, outgoing mail is STARTTLS on 587, and a failed handshake is the end of the attempt rather than a fallback to plaintext. |
| A per-service host, one host per key                                                                                                                                                | Any other API the user connects, such as Notion or Todoist. A key pasted for one service is bound to that service's host and cannot be sent anywhere else.                                                                                                                                                                                        |



## HOST

## WHAT CONNECTS TO IT

An MCP server the user connected: GitHub, Linear, Stripe, Cloudflare, or one whose address they typed themselves

New in 0.2.3, and covered in section 12. HTTPS only, port 443 only, refused at the same private-address guard as everything else, and the address is the one the user saved rather than one the assistant can name.

Arbitrary public web pages, and the provider's own search backend

Only if the user installs a skill that reads web pages or runs a research task. Private, loopback, and carrier-grade NAT addresses are refused (section 10). Web search runs on the provider's infrastructure and is billed to the user's key, so the search query reaches their search backend. Section 11 explains why that path matters more than it looks.

`assets.otianai.com`

The vendor's asset host, for large optional downloads such as the speech model or a media tool, if the user installs an add-on that needs one. Each one is checked against a published SHA-256 before it is used.

### IF YOU WANT A FIREWALL ALLOWLIST

The minimum for the app to start and stay licensed is `otianai.com`, `archie-4f35.onrender.com`, `*.googleapis.com`, and `github.com`. Everything else can be added or withheld per service, and withholding one disables that feature rather than breaking the app. Blocking the first two entirely leaves an agent on its own key working, and stops updates, asset downloads, and the periodic license check, which after roughly two months means agents refuse to start until the app can check in once.

## What stays on the computer, and what leaves it

### Never leaves the computer

- API keys, OAuth tokens, and every other credential. They stay in the operating system credential store and are read only to be put into one specific outgoing request. The key itself is then sent to the provider as a request header, which is what a key is for; what never leaves is the stored copy.
- The conversation history and the assistant's memory.
- The assistant's configuration, skills, and knowledge base files.
- Voice. Speech is turned into text on this computer and spoken replies are synthesized on this computer. No audio is uploaded.
- Anything about the local network, including any smart device the user added.
- Logs and the audit trail.

### Leaves, and to whom

- **To the user's own AI provider:** the text of the current conversation, plus whatever the assistant needed to answer, such as the body of an email it is drafting a reply to. This goes on the user's own account and key, direct from this computer.
- **To the connected service:** ordinary API calls, such as reading a message from Gmail or creating a calendar event.
- **To the chat app:** the assistant's side of the conversation, which can include the content of your mail. This is the one most reviewers underweight; it has its own subsection below.
- **To the vendor:** the account email, the plan check, and the version heartbeat. The complete list of what Otian's servers hold is below.

### THE POINT WORTH PULLING OUT, STATED WITH ITS EXCEPTIONS ATTACHED

Once an AI account of the user's own is connected, **the content of your work does not pass through a server belonging to Otian AI on its way to the model.** The computer running Archie talks to the AI provider directly, on the user's own key. That is also why the assistant stops working when the app is closed.

Three things sit beside that sentence rather than under it, and the first two are why it opens with "of the user's own".

- The **free starter credits** route AI calls through Otian's backend. Section 8 is about what that costs and what it forbids.
- **The plan sold with the AI included** routes them through the same backend for as long as the customer stays on it. It is a priced choice rather than a default, it is disclosed on the checkout screen in those words, and section 8 says exactly what it changes. **Mail is not in it:** since 3 September 2026 the assistant's email features refuse to run on this plan exactly as they do on the free credits, so no mail crosses Otian's backend on any configuration that is sold. An organization that does not want the thinking to cross it either has the plain plan, where the AI account is its own.
- **Phone access**, if turned on, is literally a queue on Otian's Firestore: messages between the computer and the phone sit there until collected. They are sealed, and the caveats on that seal are two paragraphs down.

None of the three is an edge case worth hiding in a footnote, which is why the strong sentence above is scoped rather than written as "nothing ever passes through us".

## The chat app is a real egress channel, and it is worth its own paragraph

If a user connects Telegram, Discord, Slack, Matrix, or the Messages app on a Mac, then the assistant's replies travel over that platform. When the assistant summarizes your inbox, proposes a reply to a client email, or reads back a calendar, **that content is in a message on a consumer chat platform**. For most organizations this is the largest practical exposure in the product, and it is larger than anything in the AI path, because the AI path at least runs under a commercial agreement your organization can hold.

The specifics differ by platform and they matter:

- **Telegram cloud chats are not end-to-end encrypted.** Telegram holds them and can read them. Bot conversations, which is what Archie uses, are cloud chats.
- **Slack and Discord** hold message content on their servers under their own terms. A corporate Slack you already administer is a very different proposition from a personal Discord.
- **Matrix** is the better option on this axis, because the homeserver can be one your organization runs.
- **iMessage** is different in kind from the other four. There is no bot: the assistant reads and writes in the owner's own message-to-self thread, on the Mac Archie is already running on, through the Messages app that is already signed in to their Apple account. Apple encrypts that thread between the owner's own devices. Texts from other people never reach the assistant, because the adapter is bound to that one thread and drops every other conversation before the assistant sees it. What this costs is two macOS permissions, Full Disk Access and permission to control Messages, both covered in section 14.

**Connecting a chat app is optional.** Archie's own window is a full channel on its own, and an agent that is never connected to a chat app never sends anything to one. If your review concludes that inbox content must not reach a consumer messaging platform, the configuration that satisfies you is Archie with no chat app connected, or with Matrix pointed at your own homeserver. Both are supported and neither is a degraded mode.

## Everything Otian AI's servers hold

This is the whole list, and it is the honest maximum a legal demand to Otian AI could produce. None of it is the content of anybody's work. The same list is published on the vendor's trust page, which is the canonical copy; if the two ever disagree, that one is right and this one is stale.

| WHAT                              | WHEN IT EXISTS                                                                                                                                                            |
|-----------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Your email address and account id | Always. It is how an account is identified.                                                                                                                               |
| Whether you have a current plan   | Always.                                                                                                                                                                   |
| Which add-ons you installed       | Never. Every add-on is included with Archie, and no per-person record is kept of what anybody installs; the only count is one total per add-on, with no account attached. |
| A version heartbeat               | App version, platform, edition, last seen. One record per account per edition. Can be switched off.                                                                       |
| Crash tails                       | Only if crash reporting is left on, which it can be switched off.                                                                                                         |

| WHAT                                                             | WHEN IT EXISTS                                                                                                                                                                                                                                                                                                                                                        |
|------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| The credit ledger and its spend history                          | For an account that took the free credits, and for one on the plan with the AI included: per call, the model asked for, the token counts each way, and the amount, plus salted hashes of the device and of the network address it connected from. Neither hash can be reversed into the value it was made from. Never message content, and never a prompt or a reply. |
| Second-factor records                                            | Only if two-factor sign-in is set up.                                                                                                                                                                                                                                                                                                                                 |
| Invoices, and records of a refused checkout                      | Only where a payment or a refused payment happened. A refusal record is the account id, the country, and the amount.                                                                                                                                                                                                                                                  |
| Sealed phone messages, plus their timestamps and delivery status | Only if phone access is turned on. See the caveats below.                                                                                                                                                                                                                                                                                                             |

What is not on that list is the part that matters: no conversation, no email, no calendar entry, no file, no memory, and no record of what any assistant was asked or did.

## Phone access, and the honest limits of its encryption

There is a feature for checking on an assistant from a phone at [otianai.com/phone](https://otianai.com/phone). Because the assistant lives on the one computer, the phone leaves a command in a mailbox on the vendor's Firestore and the computer picks it up. Both ends encrypt their side with AES-256-GCM under a key that this computer generates and hands to the phone by QR code. The key is never sent to Otian, so what the database holds is ciphertext Otian cannot open.

Three limits belong beside that claim, and a reviewer would find all three anyway:

- **The phone-side code is served by Otian.** The page that holds the key and does the decryption is JavaScript delivered from [otianai.com](https://otianai.com). The encryption is real, and it protects against anyone reading the database, including us. It does not protect against us changing that page. The accurate form of the claim is "we cannot read it, assuming we do not alter the page that holds the key", and anyone whose threat model includes the vendor should treat this feature as out of scope rather than as protection.
- **The metadata is not sealed.** Timestamps and delivery status are stored in the clear. When you used it and how often is visible to Otian even though what you said is not.
- **It is off by default**, per computer, and turning it off is a complete answer to all of the above.

## SECTION 8

# The two ways AI runs on Otian's account, and the gate that used to be a denylist

There are two configurations in which an AI call passes through a server belonging to Otian AI: the free credits a new install starts on, and the plan that is sold with the AI included. Every other configuration goes straight from this computer to the provider whose key the user pasted.

**The free credits.** A new install runs on two dollars of Anthropic usage paid on Otian's account, for up to fourteen days, whichever runs out first. Calls travel from this computer to Otian's backend and on to the model. The alternative was shipping Otian's own AI key inside the app, where anyone could extract it. Two dollars is a trial of Archie, not a license to it.

**The plan with the AI included.** Archie is also sold at \$59 a month or \$599 a year with \$25 of Claude Sonnet usage inside each month, so the customer opens no AI account of their own. That usage is spent on Otian's Anthropic account through the same backend, refilled on each paid month with nothing carried over, and paused when it is used up until the month rolls or the customer connects an account of their own. Two things follow that a reviewer should have in front of them:

- **Text replies do not ride it either, and that one is permanent.** The optional feature that reads a Mac's Messages database refuses on this plan for the same reason it refuses on the free credits: the people who text somebody never agreed to Otian being in the path, and the owner cannot agree on their behalf by buying a plan. Unlike the mail refusal below, no review finishing anywhere brings it back.
- **Mail does not ride it.** The refusal described below was keyed to the free credits alone until 3 September 2026, so on this plan the assistant's email features worked and the mail passed through Otian's backend on the way to Anthropic. It is now keyed to both, so no mail crosses Otian's backend on any configuration sold. The reason is not the consent argument that gates the free credits, and saying so plainly is the point of this bullet: the documents written for Google's restricted-scope verification and for the CASA assessment state that Otian operates no server in the Gmail data path, and the plan had made those statements untrue six weeks after they were written. The code was changed to match what those documents say rather than the documents amended mid-review. Mail returns to this plan when the review is finished against a submission that describes the proxy, and not before.
- **The relay keeps nothing.** The backend forwards the request and streams the answer back without writing either down, and what it records per call is the account id, the token counts each way, and the amount spent. That is a statement about what the code does today, not a cryptographic guarantee: it is a server Otian operates, so the honest claim is that nothing there logs the content, no copy is kept, and none of it is used for anything, rather than that it cannot be read.

**The configuration to approve, if either of those is a problem, is the plain plan on an AI account of your own.**

Connecting one ends the relay from that agent's next start, with nothing to switch off, and the "practical answer for an organization" in section 16 is the same answer arrived at from the provider's side.

## **On the free credits, the agent refuses to read email**

The enforcement is in the running code rather than in a policy. One arm sits ahead of every mail branch in the tool dispatch and refuses the call outright, and the background mailbox watcher refuses on the same condition. The reason is that mail content would otherwise pass through Otian's server, and the people who wrote to you did not agree to that.

THIS WAS WRONG IN AN EARLIER VERSION OF THIS DOCUMENT, AND THE CODE CHANGED RATHER THAN THE SENTENCE

Until 25 August 2026 the gate matched **tool names beginning** `inbox_` or `watch_`, and both the code and this document called that deny-by-default. It was not. A tool named `mail_fetch` or `thread_read` matches neither prefix and would have been born as exactly the hole the comment beside it promised it could not be. A naming convention is not a capability boundary, and no test can make it one, because the failure mode is somebody choosing a different name.

**It is now a classification with the default inverted.** Every tool is listed explicitly as either safe on shared credits or requiring the owner's own key, and **anything the list does not recognize is refused.** A tool wired into the dispatch tomorrow without a thought about credits is refused on the trial rather than allowed by it. A build-time test reads the declared tool names out of the source and fails if one is living on that default, so the fallback stays a backstop rather than becoming where tools quietly end up.

The rule for classifying one is written beside it: a tool is gated when running it puts somebody else's words in front of the model, which is what the proxy makes different. Not how sensitive the data feels, and not whether the tool writes.

A second layer sits at install time, over 16 add-ons that hold a kind of information the vendor has decided not to carry even in passing: the mail readers, health records and medication reminders, the add-ons that read a parent's own school mail and so end up holding details about a minor, the text-message add-on, and the personality written for a household to share. Each carries a field naming why it needs an AI account of the owner's own, and that sentence is shown when the install is declined. Two build tests enforce that the field is present and that its wording tells the reader what to do instead. This one is stricter than the mail gate above: it asks whether the agent has a key of its own, so it declines on the plan with the AI included as well as on the free credits, and each of the 16 sentences says both.

**Connecting an AI account removes the middle entirely.** Every call then goes straight from this computer to the provider on the user's own key, and every email tool works with nothing to switch on. There is a second trial for exactly that reason: fourteen days on an AI account of your own, offered when the credits have already been claimed on a computer. Nothing passes through Otian on it. **For an organization evaluating Archie, that is the trial to ask for,** because it is the same configuration you would be approving for permanent use.

## SECTION 9

# Accounts and the permissions it asks for

Mail and calendar access uses the standard OAuth flow in the user's browser, with PKCE and a validated state nonce. Archie never sees the password, and the permission screen is the provider's own.

The user ticks what they want on a chooser before the browser opens, so a person who wants calendar only never grants mailbox access. These are the scopes that can be requested:

| PROVIDER  | SCOPE                               | WHAT IT ALLOWS                                                                                                                                       |
|-----------|-------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------|
| Google    | gmail.readonly or gmail.modify      | Read mail. <code>modify</code> is requested only when the user wants the assistant to file mail out of the inbox, which is reversible.               |
|           | gmail.compose                       | Create a draft and send it. See the note below.                                                                                                      |
|           | calendar.readonly, calendar.events  | Read the calendar, and create or change events.                                                                                                      |
|           | tasks                               | Google Tasks, if the user uses to-do features.                                                                                                       |
|           | userinfo.email, drive.file          | Identify which account was connected.<br><code>drive.file</code> covers only files the user picked in the Google file picker, never the whole Drive. |
| Microsoft | Mail.Read or Mail.ReadWrite         | Read mail; <code>ReadWrite</code> for the same filing feature.                                                                                       |
|           | Mail.Send                           | Send as the user. See the note below.                                                                                                                |
|           | Calendars.Read, Calendars.ReadWrite | Read the calendar, and create or change events.                                                                                                      |
|           | User.Read, offline_access           | Identify the account, and refresh the token without asking the user to sign in again every hour.                                                     |

#### READ THIS ONE CAREFULLY, BECAUSE IT IS THE STRONGEST PERMISSION ON THE LIST

Send access is real access. What bounds it in Archie is the design of the mail features rather than the scope: the background mailbox watcher is **allowed to read and propose only**, and a drafted reply is sent when, and only when, a person presses Send on it. There is no setting that lets it send on its own, and no tool the AI can call that reaches the send path unattended.

That is a property of the software, not of the token. The token, once granted, permits sending in the way any mail client's token does, and it is held on a machine that section 2 assumes is not compromised. A reviewer who is not comfortable with that should have the user leave the send permission unticked on the consent screen.

Mail features other than sending still work.

## Mail and calendar without OAuth, on an app password

Since version 0.2.2, iCloud, Fastmail, Yahoo, AOL, Zoho and GMX connect a different way, because none of them offers Archie an OAuth flow. The user creates an **app password** at their provider and pastes it in, and Archie uses that one password for reading mail over IMAP, for sending over SMTP, and for the calendar over CalDAV. It is stored in the operating system credential store like every other credential and never written to the database.

That difference matters to a review, so it is stated rather than folded in with OAuth:

- **There is no permission chooser, because there is no consent screen to read one off.** An app password is full access to that mailbox at the provider, and Archie cannot narrow it. What Archie does instead is prove what works

at connect time and record only that: whether the mailbox opened, whether the server allows filing, whether the outgoing server accepted the same password, and whether a calendar answered.

- **Revocation is at the provider, not here.** Deleting the app password in the provider's own account settings ends Archie's access to that mailbox immediately, and it is the only thing that does.
- **The behavior above still holds.** Mail is read without marking it read, filing moves a message to Archive and never deletes it, and a reply leaves only when a person presses Send.

A further protection is worth naming, because it addresses the risk people have in mind when an AI is pointed at a mailbox. **When a model is reading an untrusted email, it is given no tools at all.** An email that tries to instruct the assistant, for example "search my mail for the password reset and quote it back", has nothing to instruct: there is no tool in the model's hands at that moment. The optional lookups a drafter may perform are named from a fixed three-item list, and every parameter of the resulting query is built by the app from facts it already holds, such as the sender's address and today's date. The model can ask for a door that exists; it cannot describe a new one.

That is the strongest anti-injection measure in the product, and it applies to one lane. Section 11 is about the lanes where the model does hold tools.

## Where Google's own verification stands, and the ceiling it puts on Gmail today

Gmail's read and send scopes are what Google calls restricted, so an application asking for them has to pass Google's verification and an independent CASA security assessment before it is treated as verified. Archie has not passed either yet, and a reviewer approving Gmail access today should know exactly where it stands rather than ask.

- **The Google submission is written and has not been sent, deliberately.** As of 3 September 2026 the app sits at Google's prepare-for-verification screen with a handful of mechanical items outstanding; the scope justifications, the demonstration video and the branding checks are done. It is being held until the CASA assessment has a scheduled date, so the two arrive together rather than the submission waiting on the assessment inside Google's queue. That is a choice with a cost, and the cost is the ceiling below.
- **The CASA assessment is engaged and paid, at assurance level 1,** and is waiting on the assessor to schedule. After that it is roughly twenty business days of lab time, then a Letter of Validation that goes to Google.
- **Until that finishes, the OAuth client is unverified, and two things follow.** Anyone connecting a Google account meets Google's own unverified-app screen and has to click through it. And an unverified client has a lifetime ceiling of **100 Google accounts**, five of which were used as of 28 August 2026. Revoking an account does not give the slot back. Otian asked Google for an exemption on 28 August 2026 and was refused the same day, so the ceiling holds until the Letter of Validation clears.

**What that means for a decision.** If your organization approves Archie for Gmail today, it is approving it on an unverified client with a finite number of connections ahead of it, and a Google refusal or delay would leave the Gmail features unavailable rather than degraded. Nothing else is affected: Outlook, the six providers that connect with an app password, calendars, and every non-Google service are outside this entirely, and so is an agent that reads no mail at all. If the timing matters to you, ask for the submission date and the assessment status in writing before you approve; they change often enough that a document is the wrong place to learn them from.



# What the AI cannot do

These are structural limits in the code, not instructions in a prompt. A prompt can be talked around. These cannot. What they do not cover is section 11.

| LIMIT                                                          | HOW IT HOLDS                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
|----------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| No shell, no code execution                                    | There is no shell anywhere: no <code>sh -c</code> , no <code>cmd /C</code> , no <code>eval</code> , and no plugin loader. Model output is never turned into a command line, and installing anything through <code>pip</code> , <code>npm</code> , or <code>brew</code> is refused outright, because those cannot be pinned or verified. <b>What is true rather than what sounds true:</b> three model-callable tools do start a program, always as a fixed argument list rather than a command string. The browser lane launches the browser already installed on the computer (the subsection below). A meeting tool passes a link the model was given to the downloader that fetches it. Media add-ons run ffmpeg or the speech model. Each of those programs is one of the hash-checked binaries in the row below, and none of them is a shell, so an argument cannot become a second command. |
| No reaching inside the network, on the paths that guard covers | Page fetching, every connected service, and every MCP call refuse any private, loopback, or carrier-grade NAT address, and refuse any hostname that resolves to one. The check runs inside the HTTP client's own name resolver, so a DNS server that answers differently the second time has no second lookup to win. <b>Two lanes sit outside that resolver and a reviewer should have both.</b> The browser lane navigates a real browser, which does its own name lookups, so an internal address is refused there only by the per-agent list of sites it may never open; the subsection below says what that means and what to do about it. And the smart-lights lane is the deliberate mirror image, allowed to reach private addresses and only those, refused loopback and carrier-grade NAT, with a test proving the two address sets cannot overlap.                                     |
| No capability it was not given                                 | The tool list is assembled fresh on every turn from that one assistant's own configuration, and it is declared in the source with typed inputs that are validated before a call runs. The one place a tool's own list is not fixed at compile time is an MCP server the user connected, whose tools are offered through a single tool that names them (section 12). There is no inherited or ambient capability, and one assistant cannot read another's memory or credentials. Every key is bound by name to the assistant that holds it, and connector keys are additionally bound to a single host. Read this as scoping rather than as containment: it bounds which tools exist in a turn, and section 11 is about what those tools can be talked into doing.                                                                                                                                 |
| Nothing that changes the outside world without a person        | Calendar changes, to-do changes, sending mail, and writes to a connected service all stage as a pending change and are shown to the user for confirmation. The confirmation is the commit. The gate covers changes that are not easily reversible; it does not cover everything, and section 11 says which.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |

| LIMIT                                           | HOW IT HOLDS                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
|-------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>No unverified binary</b>                     | Some optional add-ons download a helper program, such as ffmpeg or a speech model. An add-on can only reference a program by an opaque identifier, never by a URL, so it can point at a vetted binary and can never introduce one. The registry mapping identifier to file is the vendor's own, built from pinned sources. Every download is checked against its published SHA-256 before it is made runnable.                                                                                                                                                                                                                     |
| <b>A tamper-evident record, scoped honestly</b> | Security-relevant actions are written to an append-only log where each entry's hash chains from the one before it, and the tip is anchored in the operating system credential store and verified at startup, which detects truncation of the end of the log.<br><b>The scope of that:</b> the anchor lives in the same credential store, owned by the same user session, so anything running as that user can rewrite the log and re-anchor it. It is evidence against corruption, accident, and casual editing. It is not evidence against a local attacker, and section 2 already says a local attacker is not defended against. |

Two features go further than the rest because they touch other people's words. The mailbox watcher may read and propose only. The optional macOS text-message feature reads only forward from where it last stopped, never scans history on its own, requires Full Disk Access, which no application can grant itself, and can reach the send path through exactly one private function driven by one button a person pressed. A test counts the call sites of that function and fails the build if a second one ever appears.

## Driving a website, which is the lane that acts inside a real signed-in session

Since version 0.2.1 an assistant can do a job on a website end to end: open the site, read the page, click through it, download what it went for, and hand the file over. It is the most capable thing in the product and it is the one a review should look at hardest, so here is the whole of it.

**How it is set up.** It is off until the owner turns it on, per assistant, on the Websites panel. It drives the browser already installed on that computer, Chrome or Edge or Brave, and never a copy shipped by Otian. It uses **a profile of its own**, kept in that assistant's folder, so it does not touch the tabs, history or logins the person uses themselves. Only the owner can drive it; somebody sharing the assistant cannot.

**How a site gets signed in, and where the password goes.** The browser window opens with nothing automating it, and the person signs in themselves. Archie never sees or stores a website password: what persists afterwards is the ordinary browser session in that profile on disk. If the assistant reaches a password, card number or one-time-code field mid-job, it refuses to type into it and hands the window back to the person. The saved list of sites is a name, a host and two dates, with no credential in it, and the assistant cannot add to that list.

**What it will not do without being asked.** Ordinary clicking, reading, navigating and downloading it does on its own; that is the feature. Five things stop in code rather than in a prompt:

- A click on anything irreversible is **refused outright, rather than done once somebody has been asked**: submit, pay, delete, transfer, publish, buy, book, order, confirm, subscribe, sign up, and the phrases built out of them. The assistant's only route past one is to stop and ask, and the person clicks it themselves. There is no approval that turns this off, which is why a job shaped like filling in a form ends at the Submit button rather than through it.
- Password, card and one-time-code fields are never typed into, matched on what the page declares the field to be.

- A job has a wall-clock limit, ten minutes by default, and a cap on how many steps it may take.
- The owner can list sites the assistant may never open, checked when it opens a page, again after the page resolves, and again after every click, so a redirect into a listed site is caught rather than followed. **That list starts empty**, so it is a control to use rather than one already protecting you.
- A question to the person carries a picture of the screen and three buttons: go ahead, I will do this bit, or stop.

**Where things go.** Downloaded files land in `~/Downloads/Archie`, never in a folder the site chooses, and the site's suggested filename is scrubbed to a bare name before it is used. A picture of the screen is written to the assistant's own journal on this computer, and is sent to the person in whatever chat they are using when the assistant stops to ask something, which means it crosses that chat platform like any other message. **A screenshot is never sent to the AI model.** What the model reads is the text of the page.

**The two things to weigh.** The first is that this lane carries a live signed-in session, so a page that talks the assistant into a click is doing it inside an account somebody is already signed in to, and section 11 is where that belongs. The second is the address guard: because a real browser resolves its own names, the refusal of internal addresses described above does not cover this lane, and the control that does is the never-open list. If the computer running Archie sits on a network with internal web applications on it, that list is where they go, and the switch in section 11 turns the whole lane off in one place.

## SECTION 11

# Prompt injection, and what it can still do

---

This is the weakness. It is not solved, not by Archie and not by anybody else shipping agents today, and a document that left it out would be worth less than no document.

An AI assistant that reads a web page, a meeting transcript, a shared document, or an email is reading text that somebody else wrote. Hidden instructions in that text can try to talk the assistant into doing something the user never asked for. Everything in section 10 constrains what it can be talked into. None of it makes the assistant stop reading the text as text.

## What the gate stops

The confirmation gate stops an injected instruction from **changing anything in a connected account that is not easily reversible**. It cannot send an email on its own. It cannot delete a calendar event on its own. Every one of those stages as a proposal, and a person has to press the button.

## What the gate does not stop

An injected instruction can still, without a confirmation step:

- **Leak.** This is the serious one. The assistant can fetch a web address, and it can run a web search. Either can carry information from the conversation in the text of the request. An instruction that says "look up the following to help me" followed by content from the user's own mail produces an outbound request whose URL or query string is the leak. The SSRF guard prevents that request from reaching an internal address; it does nothing about a public one, because fetching public pages is the feature.

- **Write into memory.** The assistant can save a durable fact about its owner, and an injected instruction can put something there that persists into later conversations.
- **Add things.** Adding a to-do or a record is treated as easily reversible and is not gated the way a deletion is.
- **Act inside a website the owner is signed in to.** Where the browser lane in section 10 is switched on, the text the assistant is reading is a page it is also acting on. The irreversible clicks are refused in code, so the ceiling there is what an ordinary click can do inside somebody's signed-in account: opening, reading, navigating, downloading. That is lower than it first sounds and it is not nothing, and it is why the lane is off until somebody turns it on.

The worst realistic outcome, stated plainly: **content the assistant had access to can be sent to an attacker's server through a URL or a search query, without anybody approving anything.** That is the honest ceiling, and it should be the thing your review weighs, because it is the one Archie's architecture does not close.

## What actually reduces it

- **The no-tools rule while reading mail** (section 9) removes the channel entirely for the unattended mailbox lane, which is the highest-volume source of untrusted text most people have. This is a real mitigation and it is why that design exists.
- **There is a switch that closes the channel outright.** One setting on the agent, *Don't let this agent read the web*, withholds page fetching and the browser tools, refuses them at the dispatch if the model names them anyway, and forces every specialist's web search off however that specialist was authored. It survives installing a web skill later, which is the difference between a boundary and a configuration somebody remembers. Mail, calendar, chat, and the services connected on the Connections tab keep working.
- **Web access is off until somebody adds it.** Only a specialist can be given web search, and only if the user sets one up.
- **Fewer connected accounts is less to leak.** The assistant can only leak what it can read.

### IF THIS IS THE FINDING YOUR REVIEW STOPS ON

That is a reasonable place to stop. The configuration that reduces it furthest is the switch above, which is one checkbox you can verify rather than a list of things somebody did not install. It removes the outbound channel while leaving mail, calendar, and chat working. If your policy cannot accept even the residual paths that remain, the answer is that Archie cannot satisfy it today, and no agent product we know of can. We would rather write that than have you find it yourself in month two.

## SECTION 12

# Add-ons, and what approving Archie approves

Archie ships with a marketplace of 146 add-ons in this version. Approving Archie means approving that supply chain, so here is what it actually is.

**An add-on is data, not code.** A skill is a text file plus some settings, a personality is written instructions, a routine is a schedule. There is no mechanism by which an add-on runs as a program, so an add-on cannot itself read your disk or open a socket. That is a stronger claim than sandboxing and it is true at the operating system layer.

#### WHERE THAT CLAIM STOPS BEING REASSURING

It is true at the operating system layer and misleading if it is left there. **An add-on is instructions to an agent that holds tools**, and text that steers a tool-holding agent is capability by another name. A malicious add-on has exactly the same leak channel as any hostile text the assistant reads, which is section 11. "It cannot run code" and "it cannot cause harm" are different sentences, and only the first is true.

Two further facts a reviewer should weigh:

- **Review is two people reading submissions by hand.** Every add-on in the catalog has been read line by line by one of the two people who run Otian AI. That is a real control at this size and it is not one that scales; the vendor says so on its own trust page rather than waiting to be asked.
- **Add-on permissions are blunter than they should be.** An add-on declaring "calendar" gets reading, creating, *and deleting* together, rather than being asked about each. The cap that saves this is the account grant itself: if the user connected their calendar read-only, no add-on gets create or delete no matter what it declared. So the meaningful control is the OAuth consent screen in section 9, not the add-on's declaration.

The practical consequence for a review: **the permission decision that matters is made once, at the consent screen, not per add-on.** If you grant read-only mail and read-only calendar, nothing installed later can exceed that.

#### MCP servers, which are the one way capability arrives from outside the catalog

Version 0.2.3 added a second kind of thing a user can connect: an **MCP server**, meaning a service that publishes its own tools over the Model Context Protocol. GitHub, Linear, Stripe and Cloudflare are offered by name, and a user can type the address of any other. This is the only path by which an assistant gains tools that Otian did not write, so it belongs in a section about supply chain rather than in a list of integrations.

What bounds it:

- **Remote servers over HTTPS only.** There is no path that starts a local MCP server, because that would mean running a program, which Archie does not do anywhere. Port 443 only, and the same private-address refusal as every other outbound call, checked at connect and again on every call.
- **The address is the user's, never the model's.** The assistant gets exactly two tools: one that asks a connected server what it offers, and one that runs a named tool on a named connected server. Both take the server from a fixed list of what this user connected. There is no parameter anywhere in which the model supplies a URL.
- **Anything that changes something waits for a person.** A tool the server itself declares read-only runs on the spot. Everything else stages as a pending change and is shown for confirmation, exactly like a write to any other connected service, and a routine running unattended cannot approve one at all. The server's own claim about which of its tools are read-only decides only whether the confirmation is needed, never where the request goes or which credential rides with it.
- **The key is bound to that host.** Whatever bearer token or key the server needs is stored in the operating system credential store and cannot be sent to any other host, including across a redirect.

**Stated because this document says it will:** this feature shipped on 2026-09-02 and has not yet been exercised against a live third-party MCP server. The design above is what the code does; a first real connection has not been made.

SECTION 13

# Telemetry, and the remote switches

Archie collects a heartbeat and failure reports, and nothing else. This is the whole list, and the second row grew since the last version of this document, which is why it is now itemized.

| WHAT                           | CONTENTS                                                                                                                                                                                                                                                                                                                                                                                                                                               | HOW OFTEN                                                    |
|--------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------|
| Heartbeat                      | App version, platform, edition, and the time it was last seen. One record per account per edition.                                                                                                                                                                                                                                                                                                                                                     | At most once every six hours.                                |
| Failure reports, of four kinds | The tail of a crash and whether the app is stuck restarting; a failure to verify the audit chain at startup; a failed restore from a backup; and an add-on that has failed the same way three times, reported as the add-on's identifier and the name of the tool that failed, such as <code>skill:package-tracker</code> and <code>inbox_search</code> . Queued to disk so a crash can report itself on the next launch rather than during the crash. | Only on failure. There is no report for anything that works. |

There is no feature usage tracking, no analytics of any kind, no counts of messages or successful actions, and no message content. Every string in a report passes the same redactor that guards the logs. A desktop app is the one place a bad release cannot be pulled back, because the broken copy is on someone else's computer, so the vendor needs to know which versions are running and what is failing on them.

One transmission that is not telemetry belongs beside this list, because a reviewer watching connections will see it. Installing an add-on posts the add-on's identifier to the vendor's backend to add one to that add-on's public install count. The request is authenticated, so it is not anonymous in transit; what is stored is one running total per add-on with no account attached to it, which is what section 7's holdings list means when it says no per-person record of installs is kept.

**It can be turned off**, in Settings, under Crash reports. The same screen lists what a report contains. Turning it off stops the heartbeat, stops the uploads, and deletes anything already queued.

## The remote switches, and their authenticity story

The vendor can, without shipping a release, switch off a broken add-on, switch off a messaging channel, post a notice in the app, or set a minimum version below which agents refuse to start. This is a containment mechanism for a bad release. It cannot read anything, cannot reach a file, and cannot make the app do something it does not already do. Every default is permissive: a missing document, an absent field, and an unreachable server all behave identically, so a network that blocks the vendor entirely does not hand anyone a lever.

### THIS WAS UNSIGNED UNTIL AUGUST 2026, AND THE STATE IT IS IN NOW

The document is fetched from Firestore over HTTPS. Until 25 August 2026 its authenticity rested entirely on TLS, on Firestore's access rules, and on control of the vendor's Firebase project. Permissive defaults protect you against the document being *missing*; they never protected you against it being *wrong*, and the minimum-version field exists to stop agents starting, so a rules mistake or a compromised console was a path to stopping every install from running agents.

**The verification is now in the shipped build, which it was not when this document first said so.** The levers travel as an Ed25519-signed payload, checked against a key compiled into the build, using the same envelope Archie already uses for a signed license. A document that does not verify is ignored, which leaves the permissive defaults in place. Every release cut since 25 August 2026 carries the key, because the release script refuses to build without it, so 0.2.3 checks the signature. **Stated exactly:** a build made before that date has no key and reads the document as it always did, so for a copy still running 0.1.7 the honest description is TLS and account control.

## SECTION 14

# Permissions the operating system will ask about

Each of these produces a system prompt the user must accept. None is required for the app to run, and each corresponds to one feature.

| PROMPT           | PLATFORM           | WHAT IT IS FOR, AND WHAT HAPPENS IF IT IS DENIED                                                                                                                                                                                                                                                                                                                                                                                                           |
|------------------|--------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Microphone       | macOS and Windows  | Talking to the assistant instead of typing. Speech is turned into text on this computer and never uploaded for transcription. Denied: the microphone button does not work, and nothing else changes.                                                                                                                                                                                                                                                       |
| Local network    | macOS 15 and later | Two features on the user's own wifi: finding smart lights and plugs, and reaching an AI model the user runs on another computer of their own. Denied: the device scan finds nothing and a local model cannot be reached, and nothing else changes. On a corporate network, denying it is a reasonable default.                                                                                                                                             |
| Full Disk Access | macOS only         | Required only for the two optional features that read the Mac's Messages database: proposing replies to texts, and running the assistant in the owner's own message-to-self thread. The user must grant it by hand in System Settings; no application can grant it to itself. Denied, or simply not granted: those two features do not run, and nothing else changes. This is a broad permission and it is worth withholding unless one of them is wanted. |



| PROMPT                                                | PLATFORM   | WHAT IT IS FOR, AND WHAT HAPPENS IF IT IS DENIED                                                                                                                                                                                                                                |
|-------------------------------------------------------|------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Control Messages</b> (macOS calls this Automation) | macOS only | Asked the first time either of those two features sends a message, because sending goes through the Messages app rather than through a service of Archie's own. Denied: reading still works and sending does not, which surfaces as a plain error naming the setting to change. |

**Archie deliberately does not request Contacts access**, and one consequence of that is worth stating rather than leaving as an implication. When the text feature puts a name on a card instead of a phone number, it reads the Mac's own address book file, which sits in the same area Full Disk Access already covers. So contact names are read on a Mac where that permission was granted, without a second prompt. Adding the Contacts permission would ask the person to approve something the app does not do, which is why the key is deliberately absent.

THE MACOS ENTITLEMENTS, WHICH WERE NOT A FORMALITY

Until 25 August 2026 the macOS app was signed with two hardened-runtime relaxations, `com.apple.security.cs.disable-library-validation` and `com.apple.security.cs.allow-dyld-environment-variables`. They were there because one optional media tool is packaged with PyInstaller and unpacks its own Python libraries at launch, which library validation rejects. An earlier version of this document called that the standard entitlement set for shipping such a binary, which was true and beside the point: **library validation was off for the application process, and that process holds Keychain access.**

**They have been removed.** The app is now signed with the hardened runtime and no entitlements at all, which can be checked on any build with `codesign -d --entitlements -`. The relaxation moved to the one downloaded binary that needs it, signed separately with its own entitlements file. That binary holds no credential and reaches no Keychain, and it is run with an explicit argument list, a deadline, and a scratch directory deleted afterward.

Worth stating because it explains why this sat wrong for a while: library validation is *per process*. Once the media tools became downloaded add-ons rather than bundled sidecars, the keys on the app were doing nothing for the tool and everything to the app. The same mistake had a second half, which was that the downloaded tool was signed with no entitlements and therefore could not run at all.

SECTION 15

Signing, distribution, and updates

| PLATFORM | SIGNED             | DETAILS                                                                                                                                                                                                                                                                                                                               |
|----------|--------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| macOS    | Yes, and notarized | Signed with an Apple Developer ID certificate issued to Jett Nguyen, team ID <code>5SLXY6F2CD</code> , then submitted to Apple for notarization and stapled. Gatekeeper opens it without a warning. The signature can be checked with <code>codesign -dv --verbose=4</code> and <code>spctl -a -vv</code> against the downloaded app. |



| PLATFORM | SIGNED | DETAILS                                                                                               |
|----------|--------|-------------------------------------------------------------------------------------------------------|
| Windows  | No     | The installer is not code-signed, so SmartScreen shows a warning on first launch. See the note below. |

#### STATED PLAINLY, BECAUSE YOU WOULD FIND IT ANYWAY

The Windows installer ships unsigned. This is a gap and it is known. It is not a cost decision: the affordable route to a code-signing certificate requires three years of company history, and Otian AI LLC is not old enough yet. The practical effect is that a Windows user meets a SmartScreen warning and has to click through it, which is exactly the kind of thing an IT department is right to treat as a red flag when it appears unexplained.

What can be verified instead: the download comes from a known GitHub release under the publisher's own account, the file hash can be compared with the published release, and the macOS build of the same application is fully signed and notarized by Apple. If an unsigned installer is not acceptable under policy, the macOS build is the one to approve.

## Updates, and what the signature proves

Application updates are checked at launch against a manifest served from `otianai.com`, downloaded from GitHub, and verified against a minisign public key compiled into the app before they are applied. An update that fails that check is not installed. The update is offered to the user rather than applied silently, and requires no administrator rights.

Three things that signature does not prove, and a reviewer should have all three:

- **It does not prove the build is current.** The signature says the package came from us. An older, validly signed, vulnerable build carries a valid signature forever. Protection against a downgrade rests on the manifest, and the manifest is protected by TLS and by control of `otianai.com`, not by a signature of its own. During the beta that manifest also sits at an unlisted URL rather than behind an access check, which is obscurity and not access control.
- **The signing key is held by a two-person company.** It lives in the release environment of the person who cuts builds. There is no hardware security module, no threshold scheme, and no published rotation or revocation procedure. If that key were compromised, the recovery path today is shipping a new key in a new build, which is exactly the circular problem it sounds like.
- **The updater is the one component that cannot be fixed remotely.** A broken updater means contacting users individually, which is why the release process treats it as the highest-risk change.

## SECTION 16

# What the AI provider does with the data

Archie does not choose this. The user pastes an API key for a provider they have an account with, and that provider's terms govern what happens to the text.

Because that difference lands on whoever pasted the key, each provider's own wording is kept on file, quoted rather than paraphrased, with the page it came from and the date it was read. Twice-yearly re-reading is part of the arrangement, because a stale quote on a page whose whole worth is that it can be checked is worse than no quote. This is the current summary, read on August 16, 2026:

| PROVIDER           | TRAINS ON WHAT YOU SEND?                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            | THEIR OWN PAGE                                                                                                                                  |
|--------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------|
| Claude (Anthropic) | No                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  | <a href="https://anthropic.com/legal/commercial-terms">anthropic.com/legal/commercial-terms</a>                                                 |
| ChatGPT (OpenAI)   | No, unless you opt in. Abuse-monitoring logs are kept up to 30 days.                                                                                                                                                                                                                                                                                                                                                                                                                                                                | <a href="https://developers.openai.com/api/docs/guides/your-data">developers.openai.com/api/docs/guides/your-data</a>                           |
| Groq               | No                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  | <a href="https://console.groq.com/docs/legal/services-agreement">console.groq.com/docs/legal/services-agreement</a>                             |
| Gemini (Google)    | <b>Paid key: no. Free key: yes.</b> Google's terms say both halves under one heading: on its paid services it does not use your prompts or responses to improve its products, and on the unpaid tier it does. A key issued from a Google Cloud project with billing enabled, which is what an organization hands an employee, is the paid half. A key somebody made for themselves on the free tier is not. Archie does not ask which kind it was given, because the answer lives in that project's billing rather than in the key. | <a href="https://ai.google.dev/gemini-api/terms">ai.google.dev/gemini-api/terms</a>                                                             |
| Mistral            | <b>Free mode: yes by default.</b> Pay as you go: no.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                | <a href="https://legal.mistral.ai/terms/commercial-terms-of-service">legal.mistral.ai/terms/commercial-terms-of-service</a>                     |
| DeepSeek           | <b>Yes, and the data is stored in China,</b> if you use DeepSeek's own hosted service. Their models are open weight, so a copy running on hardware you control sends them nothing.                                                                                                                                                                                                                                                                                                                                                  | <a href="https://cdn.deepseek.com/policies/en-US/deepseek-privacy-policy.html">cdn.deepseek.com/policies/en-US/deepseek-privacy-policy.html</a> |
| Grok (xAI)         | <b>Unverified.</b> Their primary terms return an error to every automated reader, so there is no quote to stand behind. What is corroborated is that the consumer Grok app trains on conversations by default.                                                                                                                                                                                                                                                                                                                      | <a href="https://x.ai/legal">x.ai/legal</a> (unreadable as of the date above)                                                                   |

## How an API key compares with an enterprise plan

This is the question most reviewers actually want answered, and the honest answer has two halves.

**On what happens to the text, an API key gives you what an enterprise plan gives you.** Anthropic, OpenAI, and Groq all state that content sent through their API is not used to train models, and Google says the same of its paid tier, which is the tier any key issued out of a company's own Google Cloud project sits on. That is the default for an ordinary paid key rather than a term somebody has to negotiate. It is the consumer chat subscription that is the weaker product here.

**On what you can prove and enforce, it does not.** An enterprise or business agreement is a contract, and it carries things a self-service API key does not:

| WHAT ENTERPRISE ADDS               | WHY AN API KEY ALONE DOES NOT HAVE IT                                                                                                                                                           |
|------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| A signed data processing agreement | Standard API terms are accepted by clicking, not signed. If your policy requires a countersigned DPA naming your organization, that comes with the commercial agreement.                        |
| Zero data retention                | OpenAI keeps abuse-monitoring logs for up to 30 days on ordinary API use. Removing that is a separate arrangement you apply for, not a default.                                                 |
| Compliance artifacts               | SOC 2 reports, a HIPAA business associate agreement, and data residency commitments are provided under a commercial agreement, not attached to a key you generated on a web form.               |
| Administrative control             | Single sign-on, seat management, spending limits, and the ability to revoke one person's access centrally belong to the organization account. A personal key belongs to the person who made it. |

THE PRACTICAL ANSWER FOR AN ORGANIZATION

If your organization already has an agreement with an AI provider, or a Google Cloud or Google Workspace account it runs Gemini under, issue an API key from that account and have the user paste that key. Archie treats it like any other key, the traffic goes straight from the computer to your provider, and every guarantee in your existing contract applies to it without Archie being party to any of it. That is the strongest configuration available and it needs nothing special on Archie's side: the employee pastes the key their own organization issued, and the organization keeps the billing, the retention terms, and the ability to revoke it.

Running the model on your own hardware

Archie also offers a provider option called **Another provider**, which takes any endpoint speaking the widely used OpenAI-compatible format along with the model name to ask for. That covers a model your organization hosts itself. Under that arrangement the text never reaches any AI vendor: it goes from this computer to a machine you operate. For a team that cannot send its work to a third party at all, this is the route.

SECTION 17

Open points, stated rather than buried

Archie is at version 0.2.3 and is in beta. A document that lists only the good parts is not useful to somebody deciding whether to trust it, so here is the rest, including the things that have their own callouts earlier.

| POINT                                                                                        | WHAT IT MEANS FOR A REVIEWER                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |
|----------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Prompt injection can leak</b>                                                             | The largest one, and it has its own section (11). An injected instruction cannot change your accounts silently, and it can carry conversation content out through a fetched URL or a search query. Since 25 August 2026 one switch on the agent closes that channel; the residual paths section 11 names remain.                                                                                                                                                                                                                                    |
| <b>The trial mail gate was a name-prefix denylist</b>                                        | <b>Fixed 25 August 2026.</b> It is a classification with the default inverted, so an unrecognized tool is refused rather than allowed, and a build-time test keeps tools off that default. Section 8, including what the earlier version of this document got wrong.                                                                                                                                                                                                                                                                                |
| <b>The plan sold with the AI included routes the assistant's thinking through the vendor</b> | Section 8, and new since the last version of this document. On that plan the assistant's AI calls pass through Otian's backend on the way to Anthropic. Nothing there logs the content and no copy is kept, and it is a purchase decision disclosed at checkout rather than a default. Mail is excluded: email features refuse to run on this plan as they do on the free credits, so no mail crosses Otian's backend on anything sold. The configuration to approve if even the thinking is unacceptable is the plain plan on your own AI account. |
| <b>The runtime is not a separate process</b>                                                 | Section 3, and a correction: an earlier version of this document described it as a child process talking over a pipe. Archie is one process. The separation between the part that handles untrusted text and the part that can read credentials is a code boundary, not one the operating system enforces.                                                                                                                                                                                                                                          |
| <b>Driving a website acts inside a live signed-in session</b>                                | Section 10. Off until the owner turns it on, per assistant, and owner-only. Irreversible clicks are refused in code and passwords are never typed, and inside those bounds it clicks, navigates and downloads on its own. The internal-address guard does not cover it, because a real browser resolves its own names; the never-open list is the control, and it starts empty.                                                                                                                                                                     |
| <b>The weekly backup is an unencrypted copy</b>                                              | Section 4. Off by default. When it is on, it writes a plain zip of the database and every assistant's folder to a folder the person picked, and the app suggests one that syncs to cloud storage. It holds no Keychain material and it does hold conversation history.                                                                                                                                                                                                                                                                              |
| <b>MCP servers have not been exercised live</b>                                              | Section 12. The design is host-bound, guarded and gated on approval, and shipped on 2026-09-02 without a first connection to a live third-party server.                                                                                                                                                                                                                                                                                                                                                                                             |
| <b>The Windows installer is unsigned</b>                                                     | Section 15. Users meet a SmartScreen warning. The macOS build is signed and notarized.                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
| <b>Windows credential storage is weaker than macOS</b>                                       | Keychain items carry per-application access control lists. Windows Credential Manager gives user-scoped DPAPI protection with no cryptographic binding to a signed binary. Combined with an unsigned installer in a user-writable location, any process running as that user can replace the Archie executable and inherit access to its stored credentials. On Windows, treat credential protection as "protected from other users on this machine", not "protected from other software this user runs".                                           |
| <b>The macOS entitlements widened the local surface</b>                                      | <b>Fixed 25 August 2026.</b> The app is signed with no entitlements; the relaxation travels with the one downloaded tool that needs it. Section 14. Verifiable on any build with <code>codesign -d --entitlements -</code> , and not true of a build made before that date.                                                                                                                                                                                                                                                                         |

| POINT                                                                                | WHAT IT MEANS FOR A REVIEWER                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
|--------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>The remote configuration document was unsigned</b>                                | <b>Closed for current builds.</b> Section 13. Every release since 25 August 2026 carries the verification key and ignores a lever set that does not verify, and the release script will not build without it. A copy still running an older build reads the document on TLS and account control alone. Worst case there remains denial of service and a false in-app notice, not code execution.                                                                                                                                                   |
| <b>No downgrade protection, and single-person key custody</b>                        | Section 15. The signature proves origin, not recency, and the manifest that establishes recency is protected by TLS and host control alone.                                                                                                                                                                                                                                                                                                                                                                                                        |
| <b>A two-person company holds the license check</b>                                  | An installed copy runs for up to 60 days on the last license confirmation it fetched, so a company that vanished would eventually stop agents starting. What answers that is a term of the contract rather than an intention: if Otian AI ceases operations, it is obliged to publish a final version of Archie that needs no license check, within 30 days of shutting down. It is in the Terms of Service at <a href="https://otianai.com/terms-of-service">otianai.com/terms-of-service</a> , in those words, and it is there to be pointed at. |
| <b>The audit log is not attacker-proof</b>                                           | Section 10. Tamper-evident against corruption and casual editing; the anchor is reachable by the same session.                                                                                                                                                                                                                                                                                                                                                                                                                                     |
| <b>Chat platforms are a real egress path</b>                                         | Section 7. Optional, and the strongest answer is not connecting one, or using a Matrix homeserver you run.                                                                                                                                                                                                                                                                                                                                                                                                                                         |
| <b>The local database is not separately encrypted</b>                                | Conversation history sits in a SQLite file readable by anyone with access to that user's profile. No credential is in it. Whole-disk encryption is the control.                                                                                                                                                                                                                                                                                                                                                                                    |
| <b>Google has not verified the app yet, and the client has a 100-account ceiling</b> | Section 9. The submission is written and not sent, the CASA assessment is paid and waiting to be scheduled, and until both finish anyone connecting Google meets the unverified-app screen and draws on a lifetime cap of 100 accounts. Outlook, app-password mail, and every non-Google service are unaffected.                                                                                                                                                                                                                                   |
| <b>Mail scopes are real scopes</b>                                                   | Section 9. The propose-only behavior is enforced by the application; the token permits what the user granted.                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| <b>Add-on permissions are coarse</b>                                                 | Section 12. Declaring "calendar" takes read, create, and delete together; the OAuth grant is the real cap.                                                                                                                                                                                                                                                                                                                                                                                                                                         |
| <b>Some add-ons run helper binaries</b>                                              | Optional media and speech tools run as child processes with an explicit argument list, a deadline, and a scratch directory deleted afterward. They are hash-verified before they can run. They are not sandboxed by the operating system.                                                                                                                                                                                                                                                                                                          |
| <b>It is beta software and has never been audited by anyone else</b>                 | Version 0.2.3. Every feature described here works and has been exercised. There has been no external penetration test and no third-party security audit. Everything in this document is the vendor checking its own work.                                                                                                                                                                                                                                                                                                                          |

## The configuration that answers most of the table above

Every reviewer who reaches this page ends up assembling the same list out of the sections behind it, so here it is assembled. None of it is a hardening mode or a special build: each line is a setting, a permission tick, or a decision not to connect something, and each one names the section that explains it.

| DO THIS                                                                                                         | WHAT IT CLOSES                                                                                                                                                                       |
|-----------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Approve the macOS build</b>                                                                                  | Signed and notarized, with no entitlements. The Windows installer is unsigned and sits somewhere the user's own account can write. Sections 15 and 17.                               |
| <b>Have the user paste an AI key from an account your organization holds</b>                                    | Nothing then passes through Otian, your provider agreement covers the content, and you can revoke the key centrally. It also rules out both configurations in section 8. Section 16. |
| <b>Turn on whole-disk encryption</b>                                                                            | The conversation history in the local database, which is not separately encrypted. Section 4.                                                                                        |
| <b>Leave the send permission unticked at the consent screen, and grant read-only where read is all you want</b> | The gap between what the software does, which is propose only, and what the token permits. It is the one control that no later install can exceed. Sections 9 and 12.                |
| <b>Switch on "Don't let this agent read the web"</b>                                                            | The outbound channel prompt injection uses, and the browser lane with it. Mail, calendar, chat and connected services keep working. Sections 10 and 11.                              |
| <b>Connect no chat app, or a Matrix homeserver you run</b>                                                      | The largest practical egress path in the product. Archie's own window is a full channel with nothing connected. Section 7.                                                           |
| <b>If the website lane is wanted later, fill the never-open list first</b>                                      | Internal web applications, which the address guard does not cover in that lane because a real browser resolves its own names. Section 10.                                            |
| <b>Leave the weekly backup off, or point it at a folder that does not sync anywhere</b>                         | An unencrypted copy of the conversation history leaving the machine. Section 4.                                                                                                      |
| <b>Allow the four hosts in section 6 and add the rest per service</b>                                           | Everything the app does not need for the features you approved. Withholding a host disables that feature rather than breaking the app.                                               |

**What that configuration does not close, and nothing does:** the leak channel in section 11 narrows to whatever the assistant can still reach, and it does not reach zero while it reads text somebody else wrote. And it does not produce the paperwork in section 18, which is the next page and is the honest place for a review to stop if your policy needs a certificate.

### SECTION 18

## The paperwork: what exists and what does not

A review usually asks for a set of documents around hour two. Here is which of them exist, so nobody spends a week discovering the answer.

| ARTIFACT                         | EXISTS | WHERE, OR WHAT IS MISSING                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |
|----------------------------------|--------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Subprocessor list                | Yes    | Named in the privacy policy at <a href="https://otianai.com/privacy-policy">otianai.com/privacy-policy</a> : Google Firebase, Stripe, Resend, and Render. Render is the backend that handles sign-in and licensing and that relays AI requests in the two configurations in section 8, forwarding them to Anthropic on Otian's own account. In every other configuration your AI provider is not a subprocessor of ours, because you contract with them directly and the traffic never touches us.                                                                                                                                                                                                                                                                                   |
| GDPR and US state privacy rights | Yes    | The privacy policy states the lawful bases relied on in the EEA and UK (contract, legitimate interests, consent), the access, correction, deletion, portability, and objection rights, and that no personal information is sold or shared for cross-context behavioral advertising.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
| Account deletion                 | Yes    | Self-service from the account page after a recent sign-in. It removes the sign-in itself, the account record and everything filed under it, including invoices and any second-factor records, and it deletes the customer object and stored payment methods at Stripe. Three things outlive it because they are keyed separately: failure reports, the version heartbeat, and the credit ledger with its spend history and its device and address hashes. None of those contains anything the user wrote, and they can be deleted on request.                                                                                                                                                                                                                                        |
| Retention schedule with numbers  | Yes    | A period per record, in the privacy policy at <a href="https://otianai.com/privacy-policy#data-retention">otianai.com/privacy-policy#data-retention</a> : 90 days for failure reports and daily address hashes, 12 months for spend history, device hashes and refusal records, the life of the account plus 90 days for the credit ledger, 7 years for invoices because company records law says so, and 30 days for an uncollected phone message. Each record now carries its own deletion date and the database acts on it. Two things are said beside the table rather than left out of it: our database provider keeps its own backups for seven days, so a deletion is a deletion from the live systems first, and deleting an account does not wait for any of these periods. |
| Vulnerability disclosure policy  | Yes    | At <a href="https://otianai.com/security">otianai.com/security</a> , with <a href="https://.well-known/security.txt">/.well-known/security.txt</a> pointing at it. It carries the times we hold ourselves to: a reply within one working day, and within five, whether we agree it is a problem and when it will be fixed. Reports go to <a href="mailto:questions@otianai.com">questions@otianai.com</a> , which is the address on this document rather than a separate security alias, because it is the mailbox that is actually read. There is no advisory feed and no bounty program, and the page says both.                                                                                                                                                                   |

| ARTIFACT                                           | EXISTS               | WHERE, OR WHAT IS MISSING                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |
|----------------------------------------------------|----------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Incident and breach notification commitment</b> | <b>Yes</b>           | <b>72 hours</b> from establishing that personal information we hold about you was exposed, by email to the account address, saying what was taken, when, what we have done, and what you should do, including where the honest answer is that we do not know yet. It is in the privacy policy and it is a clause of the Terms of Service, beside the one about ceasing operations, so it is a term rather than a page. What limits the blast radius is still the holdings list in section 7: a breach of Otian AI cannot produce anybody's conversations, mail, calendar, files, or keys, because none of those are held. |
| <b>Data processing agreement</b>                   | <b>On request</b>    | Send us yours and we will review and sign it. There is no published Otian standard agreement yet, and the security page says so rather than implying one exists. What it would cover is short, because on your own AI account the content never reaches us: the account email, the plan status, and the operational records in section 7.                                                                                                                                                                                                                                                                                 |
| <b>SOC 2, ISO 27001, HIPAA BAA</b>                 | <b>No</b>            | None held, and none in progress. A two-person company at version 0.2.3. If your policy requires a certification from the software vendor, Archie does not meet it today, and no amount of architecture in this document changes that answer.                                                                                                                                                                                                                                                                                                                                                                              |
| <b>Dependency inventory and CVE handling</b>       | <b>Not published</b> | Dependencies are pinned in checked-in lock files for both the Rust and TypeScript sides, and four packages are deliberately held below latest with the blocker recorded for each. There is no published SBOM and no stated CVE response window. An SBOM can be produced on request from the lock files.                                                                                                                                                                                                                                                                                                                   |
| <b>External security audit</b>                     | <b>No</b>            | Nobody has audited Archie but the people who wrote it.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    |

## Which document this is

This one is revision 5. The number exists because a reviewer pointed out that two different versions of this document had gone out carrying the same product version and the same date, so "Archie 0.2.3, prepared 3 September 2026" no longer named one document. If you are recording what you approved, record the revision number with the build.

| REVISION                     | WHAT CHANGED                                                                                                                                                                                                                                                                                                                                                                              |
|------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>1</b><br>25 August 2026   | First issue, describing version 0.1.7.                                                                                                                                                                                                                                                                                                                                                    |
| <b>2</b><br>3 September 2026 | Brought to version 0.2.3. Corrected four claims that had become wrong: the runtime is not a separate process, the log and file paths, the remote-lever signing, and the AI-included plan as a second path through the vendor. Added the browser lane, MCP servers, app-password mail, the weekly backup, the Automation permission, and the configuration table at the end of section 17. |
| <b>3</b><br>3 September 2026 | Mail came off the AI-included plan in the code, and section 8 says why. Added where Google's own verification stands and the 100-account ceiling that comes with it (section 9), and this table.                                                                                                                                                                                          |



| REVISION              | WHAT CHANGED                                                                                                                                                                                                                                                                                                                                                       |
|-----------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 4<br>3 September 2026 | Section 9 says the submission is being held until the assessment is scheduled, rather than leaving "not sent" to be read as inattention. Section 8's install gate covers 16 add-ons and said 19, and each of those 16 now names this plan as well as the free credits in its own refusal.                                                                          |
| 5<br>3 September 2026 | Four rows in section 18 turned from No to Yes, because the things they described as missing now exist: a published disclosure policy, a retention schedule with a period per record, a 72-hour breach commitment, and an answer on data processing agreements. Text replies came off the plan with the AI included, for a reason that does not expire (section 8). |

#### REMOVING IT

Archie is a normal application and uninstalls like one. Deleting the app removes the program; deleting the folder named in section 4 removes its data; the credentials it created appear under `com.archie.app` in Keychain Access or Windows Credential Manager and can be deleted there. Two things sit outside all of that and have to be dealt with by hand: files the assistant saved, in `~/Downloads/Archie`, and, if weekly backups were switched on, the folder chosen for them, which holds whole copies of the database. Access already granted to a mailbox or calendar should be revoked at the provider, in the Google account permissions page or the Microsoft 365 admin center, which is where OAuth grants actually live.

**Questions.** Anything in this document can be checked further, including specific network captures, the signing details on a given build, or a walkthrough of a feature before it is approved. If your review stopped on one row, say which one: we would rather answer that than send more pages. Write to Jack at [jack@otianai.com](mailto:jack@otianai.com).

Archie 0.2.3 · Revision 5, September 3, 2026 · Otian AI LLC · [otianai.com](https://otianai.com). Details describe this version and may change in later releases.