

# What Archie is, and what it does on the computer it runs on

---

A plain description of the software for anyone who has been asked to approve it: where it installs, what it sends over the network, where credentials are kept, and what it is not allowed to do.

---

## PRODUCT

**Archie, version 0.1.7 (beta)**

## PUBLISHER

**Otian AI LLC**

## PLATFORMS

**macOS and Windows**

## DOCUMENT DATE

**August 25, 2026**

## CONTACT

**jett@otianai.com  
jack@otianai.com**

## WEBSITE

**otianai.com**

## SECTION 1

# What Archie is

---

Archie is a desktop application that lets one person set up an AI assistant and then talk to it from a chat app they already use.

It is installed like any other desktop program and it runs only while it is open. There is no Archie server that holds the assistant: the assistant is a process on the computer where the app is installed. Close the app and it stops answering.

The person using it does three things. First they connect an **AI account**. This is an API key from a provider such as Anthropic, OpenAI, or Google: a developer account billed by what it is used for, not a consumer chat subscription like ChatGPT Plus or Claude Pro. The difference matters for a review and section 16 sets it out, but the short form is that the major providers do not train on anything sent through their API, and that is the default rather than a setting somebody has to find. Second, they optionally connect a chat app, such as Telegram, Slack, Discord, Signal, or Matrix, so they can reach the assistant from their phone in addition to their computer. Third, they install **skills**, which are small bundles of instructions that tell the assistant how to handle one kind of job, like drafting a reply to an email or putting an event on a calendar.

What the assistant can actually touch is decided entirely by what the user connects. A fresh install can do nothing except talk. Reading a mailbox requires that person to sign in to that mailbox and grant permission. Reading a calendar is a separate permission. Nothing is on by default and nothing is inherited from another user.

### THE SHORT VERSION FOR A REVIEWER

Archie is a locally installed client. It makes outbound HTTPS connections to the services the user connected and to the vendor's own backend. It opens no listening network port, accepts no inbound connection, runs no shell command, and installs no service, driver, or kernel extension. It does not require administrator rights to install on Windows.

It is also an AI agent that reads text other people wrote, and that has a class of weakness no vendor has solved. Section 11 is about that, and it is the section to read first if you only read one.

## SECTION 2

# What this document claims, and what it does not

---

A security claim means nothing without the boundary it holds inside. This section is that boundary, stated before anything else, so every sentence after it can be checked against something.

## What Archie is built to defend against

- **A remote attacker on the network.** There is nothing to connect to. No listening port, no inbound path, no service.
- **Model output being turned into execution.** No shell, no eval, no dynamic code loading, anywhere.
- **An assistant being talked into reaching your internal network.** Private, loopback, and carrier-grade NAT addresses are refused, and the refusal happens inside the name resolver.

- **Secrets sitting in readable places.** Credentials live in the operating system credential store and are read transiently. They are never in the database, in logs, or in exports.
- **A malicious or broken add-on introducing a binary.** An add-on can reference a vetted, hash-pinned binary by identifier. It can never name a URL.
- **An action being taken without a person.** Sending mail, changing a calendar, and writing to a connected service all stage as a proposal and commit on a human action.

What it does not defend against, stated plainly

| NOT DEFENDED                                   | WHAT THAT MEANS FOR YOU                                                                                                                                                                                                           |
|------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| An endpoint that is already compromised        | Archie assumes the computer it runs on is trusted. Anything already running as that user can read Archie's database, its files, and, on Windows, its credentials. Nothing in this document survives a compromised machine.        |
| A local attacker inside the user's own session | The audit log, the credential store, and the update path are all reachable by anything running as that user. See sections 10, 15, and 17.                                                                                         |
| The user themselves                            | Archie is the user's tool. It does not stop them from connecting an account, approving a draft, or pasting a key. There is no administrator above them in the personal edition, and the license gates are on a computer they own. |
| Prompt injection, fully                        | The gate stops an injected instruction from changing things silently. It does not stop it from leaking. This is the real one, and section 11 is the whole story rather than a line here.                                          |
| What your AI provider does                     | Content goes to the provider whose key was pasted, under your agreement with them. Archie is not party to it. Section 16.                                                                                                         |
| What a chat platform does                      | If a chat app is connected, the assistant's messages transit that platform on its terms. Section 7 covers what that actually exposes.                                                                                             |

ONE RULE THIS DOCUMENT TRIES TO FOLLOW

Where it says "everything" or "the whole list", it means it, and you should check it. Where it cannot be complete it says so instead. If you find something Archie does that is not described here, that is a defect in this document and worth telling us about, because the value of a page like this is entirely in whether it survives being checked.

SECTION 3

How the application is built

Archie is three parts, each allowed to do less than the one below it. The split exists so that the part handling untrusted text is the part furthest from the credentials.

### The window

NO SECRETS

The screens people look at, written in React and TypeScript and drawn by the operating system's own web view. It holds no credentials, runs no database queries, and cannot read or write files. It asks the layer below for everything, through a fixed list of named commands that is deny-by-default.

↓ **typed commands**, allowlisted one by one

### The core

HOLDS THE KEYS

Rust. Owns the local database, the operating system credential store, and the audit log. This is the only part that can read a secret value, and it reads one only at the moment it is putting it into a specific outgoing request. A secret is never handed back up to the window after it is typed in.

↓ **standard input and output only**, no network port

### The runtime

RUNS THE ASSISTANT

A separate child process that does the actual work: deciding which skill a message belongs to, calling the AI provider, and running tools. It is a separate process so that a hang or a crash can be restarted without taking the app down with it. It talks to the core over the pipe it was started with, not over the network.

The application is written in Rust and TypeScript and packaged with Tauri 2, which means it uses the operating system's built-in web view rather than shipping a copy of a browser.

**One caveat about that word deny-by-default.** It is accurate for the boundary above, where the window may only call commands that were explicitly granted. It is not true of every boundary inside Archie, and section 8 names one place where a gate is written as a denylist instead. The two are described differently in this document on purpose.

## SECTION 4

# What gets installed, and where

| PLATFORM | INSTALLER                                                                            | RIGHTS NEEDED                                                                                                                      |
|----------|--------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------|
| macOS    | A signed disk image ( <code>.dmg</code> ). The user drags the app into Applications. | None beyond the ability to write to Applications. It can also be run from the user's own folder.                                   |
| Windows  | An NSIS installer ( <code>.exe</code> ), built in <b>current user</b> mode.          | None. It installs under the signed-in user's own profile and does not write to Program Files or to the machine-wide registry hive. |

No background service, scheduled task, daemon, driver, or kernel extension is installed. Nothing runs at startup unless the user turns that on. When the app is closed, no Archie process remains.

**The Windows install location is a tradeoff, not a free win.** Needing no administrator rights is convenient and it also means the installed executable sits somewhere the user's own account can write. Section 17 says what follows from that.

## Files it creates

| WHAT                                                | WHERE ON MACOS                                                      | HOW IT IS STORED                                                                    |
|-----------------------------------------------------|---------------------------------------------------------------------|-------------------------------------------------------------------------------------|
| Settings, agent configuration, conversation history | <code>~/Library/Application Support/com.archie.app/archie.db</code> | A SQLite file. Not separately encrypted. See the note below.                        |
| Credentials and API keys                            | macOS Keychain, service <code>com.archie.app</code>                 | Operating system credential store. Never written to the database or to any file.    |
| Logs                                                | <code>~/Library/Logs/Archie/</code>                                 | Rotating JSON lines. Every line passes a redactor that strips secret-shaped values. |
| Caches and files the assistant saved                | <code>~/Library/Caches/com.archie.app/</code>                       | Ordinary files.                                                                     |

On Windows the same four live under the user's `AppData` folder, and credentials live in Windows Credential Manager instead of the Keychain. Those two stores are not equivalent, and section 17 says how.

### WORTH KNOWING

The local database is a plain SQLite file. It holds settings, the assistant's configuration, and conversation history, and anyone who can read that user's profile folder can read it. It never holds an API key, a password, or an OAuth token. If conversation history on the machine is a concern, the control is whole-disk encryption: FileVault on macOS or BitLocker on Windows.

## SECTION 5

# Network behavior

Every connection Archie makes is outbound HTTPS or a secure websocket that it opened itself. Nothing on the network can connect to it.

## Inbound: none

Archie listens on no port and accepts no inbound connection. There is one narrow exception, and it lasts seconds:

- When the user signs in to Google, Microsoft, or their Otian account, the app opens a listener on `127.0.0.1` on a port the operating system picks, because that is how the standard OAuth desktop flow returns the user to the app after they approve in their browser. It is bound to the loopback address, so it is not reachable from the network, and it closes as soon as the browser hands back the response.
- That flow uses PKCE** (`archie-core::oauth` generates a verifier and challenge per attempt) together with a state nonce that is validated on return. Both matter here: on a shared machine, a loopback redirect without PKCE is interceptable by any other local process, and PKCE is what makes an intercepted authorization code useless.

The runtime that executes the assistant has no socket of its own at all. It communicates with the rest of the app over standard input and output.

Local network: only if the user asks for it

Archie has an optional feature for switching smart lights and plugs on the user's own wifi. It is off until the user runs a scan and adds a device by hand. If it is used, the app sends UDP discovery broadcasts on the local subnet. On a corporate network this feature has nothing to find and can simply be left alone; there is no need to use it.

SECTION 6

Where it connects, and why

This is the full list of destinations, grouped by what causes the connection. Anything in the second group happens only because the user connected that service.

Always, as part of the application working

| HOST                                | PURPOSE                                 | NOTES                                                                                                                                                                                                                                                                                |
|-------------------------------------|-----------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| otianai.com                         | Sign-in, license check, update manifest | The vendor's own site. Also serves the optional phone page.                                                                                                                                                                                                                          |
| archie-4f35.onrender.com            | The vendor's backend service            | Sign-in, licensing, account management, account email, and, while a user is on the free starter credits, relaying AI requests to Anthropic. This is the host to add to a block list alongside otianai.com if you want to test that an agent on its own key keeps working without us. |
| *.googleapis.com                    | Firebase Auth and Firestore             | The vendor's Firebase project holds the account record, the add-on catalog, the remote configuration document, and the license tier. Installing a free add-on also increments a global counter on that add-on's catalog document; no per-person record of free installs is kept.     |
| github.com                          | Downloading an app update               | Update binaries are GitHub release assets. Every update package is signature-verified before it is applied. Section 15 says what that does and does not prove.                                                                                                                       |
| api.stripe.com, checkout.stripe.com | Subscription and billing                | Opened in the browser. Archie never sees a card number.                                                                                                                                                                                                                              |

Only after the user connects that service

| HOST                                                                                            | WHAT CONNECTS TO IT                                                                           |
|-------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------|
| api.anthropic.com, api.openai.com, generativelanguage.googleapis.com, api.groq.com, and similar | The AI provider whose key the user pasted in. This is where the assistant's thinking happens. |

| HOST                                                                                                                                            | WHAT CONNECTS TO IT                                                                                                                                                                                                                                                                                                                                            |
|-------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>gmail.googleapis.com</code> , <code>www.googleapis.com</code> ,<br><code>tasks.googleapis.com</code> , <code>oauth2.googleapis.com</code> | Gmail, Google Calendar, and Google Tasks, if the user signs in to a Google account.                                                                                                                                                                                                                                                                            |
| <code>graph.microsoft.com</code> , <code>login.microsoftonline.com</code>                                                                       | Outlook mail and calendar, if the user signs in to a Microsoft 365 account.                                                                                                                                                                                                                                                                                    |
| <code>api.telegram.org</code> , <code>discord.com</code> , <code>slack.com</code> , the user's Matrix homeserver, Signal's servers              | Whichever chat app the user connected. Read section 7 before approving this one.                                                                                                                                                                                                                                                                               |
| A per-service host, one host per key                                                                                                            | Any other API the user connects, such as Notion or Todoist. A key pasted for one service is bound to that service's host and cannot be sent anywhere else.                                                                                                                                                                                                     |
| Arbitrary public web pages, and the provider's own search backend                                                                               | Only if the user installs a skill that reads web pages or runs a research task. Private, loopback, and carrier-grade NAT addresses are refused (section 10). Web search runs on the provider's infrastructure and is billed to the user's key, so the search query reaches their search backend. Section 11 explains why that path matters more than it looks. |
| The vendor's asset host                                                                                                                         | Large optional downloads, such as the speech model or a media tool, if the user installs an add-on that needs one. Each one is checked against a published SHA-256 before it is used.                                                                                                                                                                          |

#### IF YOU WANT A FIREWALL ALLOWLIST

The minimum for the app to start and stay licensed is `otianai.com`, `archie-4f35.onrender.com`, `*.googleapis.com`, and `github.com`. Everything else can be added or withheld per service, and withholding one disables that feature rather than breaking the app. Blocking the first two entirely leaves an agent on its own key working, and stops updates, asset downloads, and the periodic license check, which after roughly two months means agents refuse to start until the app can check in once.

## What stays on the computer, and what leaves it

### Never leaves the computer

- API keys, OAuth tokens, and every other credential. They stay in the operating system credential store and are read only to be put into one specific outgoing request. The key itself is then sent to the provider as a request header, which is what a key is for; what never leaves is the stored copy.
- The conversation history and the assistant's memory.
- The assistant's configuration, skills, and knowledge base files.
- Voice. Speech is turned into text on this computer and spoken replies are synthesized on this computer. No audio is uploaded.
- Anything about the local network, including any smart device the user added.
- Logs and the audit trail.

### Leaves, and to whom

- **To the user's own AI provider:** the text of the current conversation, plus whatever the assistant needed to answer, such as the body of an email it is drafting a reply to. This goes on the user's own account and key, direct from this computer.
- **To the connected service:** ordinary API calls, such as reading a message from Gmail or creating a calendar event.
- **To the chat app:** the assistant's side of the conversation, which can include the content of your mail. This is the one most reviewers underweight; it has its own subsection below.
- **To the vendor:** the account email, the plan check, and the version heartbeat. The complete list of what Otian's servers hold is below.

### THE POINT WORTH PULLING OUT, STATED WITH ITS EXCEPTIONS ATTACHED

Once an AI account is connected, **the content of your work does not pass through a server belonging to Otian AI on its way to the model.** The computer running Archie talks to the AI provider directly, on the user's own key. That is also why the assistant stops working when the app is closed.

Two things sit beside that sentence rather than under it. The **free starter credits** route AI calls through Otian's backend, and section 8 is about what that costs and what it forbids. **Phone access**, if turned on, is literally a queue on Otian's Firestore: messages between the computer and the phone sit there until collected. They are sealed, and the caveats on that seal are two paragraphs down. Neither is an edge case worth hiding in a footnote, which is why the strong sentence above is scoped to the AI path rather than written as "nothing ever passes through us".

## The chat app is a real egress channel, and it is worth its own paragraph

If a user connects Telegram, Discord, Slack, Signal, or Matrix, then the assistant's replies travel over that platform. When the assistant summarizes your inbox, proposes a reply to a client email, or reads back a calendar, **that content is in a message on a consumer chat platform.** For most organizations this is the largest practical exposure in the product, and it is larger than anything in the AI path, because the AI path at least runs under a commercial agreement your organization can hold.

The specifics differ by platform and they matter:

- **Telegram cloud chats are not end-to-end encrypted.** Telegram holds them and can read them. Bot conversations, which is what Archie uses, are cloud chats.



- **Slack and Discord** hold message content on their servers under their own terms. A corporate Slack you already administer is a very different proposition from a personal Discord.
- **Signal and Matrix** are the better options on this axis, and Matrix can be a homeserver your organization runs.

**Connecting a chat app is optional.** Archie's own window is a full channel on its own, and an agent that is never connected to a chat app never sends anything to one. If your review concludes that inbox content must not reach a consumer messaging platform, the configuration that satisfies you is Archie with no chat app connected, or with Matrix pointed at your own homeserver. Both are supported and neither is a degraded mode.

## Everything Otian AI's servers hold

This is the whole list, and it is the honest maximum a legal demand to Otian AI could produce. None of it is the content of anybody's work. The same list is published on the vendor's trust page, which is the canonical copy; if the two ever disagree, that one is right and this one is stale.

| WHAT                                                             | WHEN IT EXISTS                                                                                                                                                                            |
|------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Your email address and account id                                | Always. It is how an account is identified.                                                                                                                                               |
| Whether you have a current plan                                  | Always.                                                                                                                                                                                   |
| Which <b>paid</b> add-ons you bought                             | Per purchase, kept permanently as a receipt so a paid add-on can be reinstalled on another computer. No per-person record is kept of the free add-ons anybody installs.                   |
| A version heartbeat                                              | App version, platform, edition, last seen. One record per account per edition. Can be switched off.                                                                                       |
| Crash tails                                                      | Only if crash reporting is left on, which it can be switched off.                                                                                                                         |
| The trial-credit ledger and its spend history                    | Only for an account that took the free credits: token counts and amounts, with salted hashes of the device and the IP address so the grant limits can be enforced. Never message content. |
| Second-factor records                                            | Only if two-factor sign-in is set up.                                                                                                                                                     |
| Invoices, and records of a refused checkout                      | Only where a payment or a refused payment happened. A refusal record is the account id, the country, and the amount.                                                                      |
| Sealed phone messages, plus their timestamps and delivery status | Only if phone access is turned on. See the caveats below.                                                                                                                                 |

What is not on that list is the part that matters: no conversation, no email, no calendar entry, no file, no memory, and no record of what any assistant was asked or did.

## Phone access, and the honest limits of its encryption

There is a feature for checking on an assistant from a phone at [otianai.com/phone](https://otianai.com/phone). Because the assistant lives on the one computer, the phone leaves a command in a mailbox on the vendor's Firestore and the computer picks it up. Both ends encrypt their side with AES-256-GCM under a key that this computer generates and hands to the phone by QR code. The key is never sent to Otian, so what the database holds is ciphertext Otian cannot open.

Three limits belong beside that claim, and a reviewer would find all three anyway:

- **The phone-side code is served by Otian.** The page that holds the key and does the decryption is JavaScript delivered from `otianai.com`. The encryption is real, and it protects against anyone reading the database, including us. It does not protect against us changing that page. The accurate form of the claim is "we cannot read it, assuming we do not alter the page that holds the key", and anyone whose threat model includes the vendor should treat this feature as out of scope rather than as protection.
- **The metadata is not sealed.** Timestamps and delivery status are stored in the clear. When you used it and how often is visible to Otian even though what you said is not.
- **It is off by default**, per computer, and turning it off is a complete answer to all of the above.

## SECTION 8

# The free trial, and the one gate written as a denylist

A new install runs on two dollars of Anthropic usage paid on Otian's account, for up to fourteen days. It is the one configuration in which an AI call passes through a server belonging to Otian AI.

During the trial, calls travel from this computer to Otian's backend and on to the model. The alternative was shipping Otian's own AI key inside the app, where anyone could extract it. Two dollars is a trial of Archie, not a license to it.

## On the trial credits, the agent refuses to read email

The enforcement is in the running code rather than in a policy. One arm sits ahead of every mail branch in the tool dispatch and refuses the call outright, and the background mailbox watcher refuses on the same condition. The reason is that mail content would otherwise pass through Otian's server, and the people who wrote to you did not agree to that.

### THE HONEST SHAPE OF THAT GATE, BECAUSE IT IS NOT WHAT SECTION 3 DESCRIBES

That arm matches on **tool names beginning** `inbox_` or `watch_`. It is a denylist on a string prefix, not a capability boundary, and those are not the same thing. A future tool named `mail_fetch` or `thread_read` would not be caught by it. The naming convention is enforced by convention and by tests, and a test asserts that names follow the pattern; it cannot assert that a capability was declared, because no capability tag exists to declare.

An earlier draft of this document said such a tool would be "born gated rather than born as a hole". That was the intent behind the design and it was not an accurate description of the mechanism, so it has been removed rather than softened. The correct statement is: **every mail tool that exists today is gated, and the gate depends on a naming convention holding**. Making it depend on a declared capability instead is on the list in section 17.

A second layer sits at install time over the add-ons that hold a legally distinct kind of information: health records, medication reminders, the text-message add-on, and the two school add-ons, which read a parent's own school mail and so end up holding details about a minor. Each carries a field naming why it needs an AI account of the owner's

own, and that sentence is shown when the install is declined. Two build tests enforce that the field is present and that its wording tells the reader what to do instead.

**Connecting an AI account removes the middle entirely.** Every call then goes straight from this computer to the provider on the user's own key, and every email tool works with nothing to switch on. There is a second trial for exactly that reason: fourteen days on an AI account of your own, offered when the credits have already been claimed on a computer. Nothing passes through Otian on it. **For an organization evaluating Archie, that is the trial to ask for**, because it is the same configuration you would be approving for permanent use.

SECTION 9

## Accounts and the permissions it asks for

Mail and calendar access uses the standard OAuth flow in the user's browser, with PKCE and a validated state nonce. Archie never sees the password, and the permission screen is the provider's own.

The user ticks what they want on a chooser before the browser opens, so a person who wants calendar only never grants mailbox access. These are the scopes that can be requested:

| PROVIDER  | SCOPE                               | WHAT IT ALLOWS                                                                                                                                       |
|-----------|-------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------|
| Google    | gmail.readonly or gmail.modify      | Read mail. <code>modify</code> is requested only when the user wants the assistant to file mail out of the inbox, which is reversible.               |
|           | gmail.compose                       | Create a draft and send it. See the note below.                                                                                                      |
|           | calendar.readonly, calendar.events  | Read the calendar, and create or change events.                                                                                                      |
|           | tasks                               | Google Tasks, if the user uses to-do features.                                                                                                       |
|           | userinfo.email, drive.file          | Identify which account was connected.<br><code>drive.file</code> covers only files the user picked in the Google file picker, never the whole Drive. |
| Microsoft | Mail.Read or Mail.ReadWrite         | Read mail; <code>ReadWrite</code> for the same filing feature.                                                                                       |
|           | Mail.Send                           | Send as the user. See the note below.                                                                                                                |
|           | Calendars.Read, Calendars.ReadWrite | Read the calendar, and create or change events.                                                                                                      |
|           | User.Read, offline_access           | Identify the account, and refresh the token without asking the user to sign in again every hour.                                                     |

READ THIS ONE CAREFULLY, BECAUSE IT IS THE STRONGEST PERMISSION ON THE LIST

Send access is real access. What bounds it in Archie is the design of the mail features rather than the scope: the background mailbox watcher is **allowed to read and propose only**, and a drafted reply is sent when, and only when, a person presses Send on it. There is no setting that lets it send on its own, and no tool the AI can call that reaches the send path unattended.

That is a property of the software, not of the token. The token, once granted, permits sending in the way any mail client's token does, and it is held on a machine that section 2 assumes is not compromised. A reviewer who is not comfortable with that should have the user leave the send permission unticked on the consent screen. Mail features other than sending still work.

A further protection is worth naming, because it addresses the risk people have in mind when an AI is pointed at a mailbox. **When a model is reading an untrusted email, it is given no tools at all.** An email that tries to instruct the assistant, for example "search my mail for the password reset and quote it back", has nothing to instruct: there is no tool in the model's hands at that moment. The optional lookups a drafter may perform are named from a fixed three-item list, and every parameter of the resulting query is built by the app from facts it already holds, such as the sender's address and today's date. The model can ask for a door that exists; it cannot describe a new one.

That is the strongest anti-injection measure in the product, and it applies to one lane. Section 11 is about the lanes where the model does hold tools.

SECTION 10

## What the AI cannot do

These are structural limits in the code, not instructions in a prompt. A prompt can be talked around. These cannot. What they do not cover is section 11.

| LIMIT                          | HOW IT HOLDS                                                                                                                                                                                                                                                                                                                                                                                                         |
|--------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| No shell, no code execution    | There is no tool that runs a command, evaluates code, or loads a plugin. Tools are a fixed list declared in the source with typed inputs that are validated before the call runs. Model output is never turned into a command line. Installing anything through <code>pip</code> , <code>npm</code> , or <code>brew</code> is refused outright, because those cannot be pinned or verified.                          |
| No reaching inside the network | The web tools refuse any private, loopback, or carrier-grade NAT address, and refuse any hostname that resolves to one. The check runs inside the HTTP client's own name resolver, so a DNS server that answers differently the second time has no second lookup to win. This is a standing defense against server-side request forgery, and it means the assistant cannot be talked into fetching an internal page. |

| LIMIT                                                   | HOW IT HOLDS                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
|---------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| No capability it was not given                          | The tool list is assembled fresh on every turn from that one assistant's own configuration. There is no inherited or ambient capability, and one assistant cannot read another's memory or credentials. Every key is bound by name to the assistant that holds it, and connector keys are additionally bound to a single host. Read this as scoping rather than as containment: it bounds which tools exist in a turn, and section 11 is about what those tools can be talked into doing.                                                                                                                                          |
| Nothing that changes the outside world without a person | Calendar changes, to-do changes, sending mail, and writes to a connected service all stage as a pending change and are shown to the user for confirmation. The confirmation is the commit. The gate covers changes that are not easily reversible; it does not cover everything, and section 11 says which.                                                                                                                                                                                                                                                                                                                        |
| No unverified binary                                    | Some optional add-ons download a helper program, such as ffmpeg or a speech model. An add-on can only reference a program by an opaque identifier, never by a URL, so it can point at a vetted binary and can never introduce one. The registry mapping identifier to file is the vendor's own, built from pinned sources. Every download is checked against its published SHA-256 before it is made runnable.                                                                                                                                                                                                                     |
| A tamper-evident record, scoped honestly                | Security-relevant actions are written to an append-only log where each entry's hash chains from the one before it, and the tip is anchored in the operating system credential store and verified at startup, which detects truncation of the end of the log.<br><b>The scope of that:</b> the anchor lives in the same credential store, owned by the same user session, so anything running as that user can rewrite the log and re-anchor it. It is evidence against corruption, accident, and casual editing. It is not evidence against a local attacker, and section 2 already says a local attacker is not defended against. |

Two features go further than the rest because they touch other people's words. The mailbox watcher may read and propose only. The optional macOS text-message feature reads only forward from where it last stopped, never scans history on its own, requires Full Disk Access, which no application can grant itself, and can reach the send path through exactly one private function driven by one button a person pressed. A test counts the call sites of that function and fails the build if a second one ever appears.

SECTION 11

## Prompt injection, and what it can still do

This is the weakness. It is not solved, not by Archie and not by anybody else shipping agents today, and a document that left it out would be worth less than no document.

An AI assistant that reads a web page, a meeting transcript, a shared document, or an email is reading text that somebody else wrote. Hidden instructions in that text can try to talk the assistant into doing something the user never asked for. Everything in section 10 constrains what it can be talked into. None of it makes the assistant stop reading the text as text.

## What the gate stops

The confirmation gate stops an injected instruction from **changing anything in a connected account that is not easily reversible**. It cannot send an email on its own. It cannot delete a calendar event on its own. Every one of those stages as a proposal, and a person has to press the button.

## What the gate does not stop

An injected instruction can still, without a confirmation step:

- **Leak.** This is the serious one. The assistant can fetch a web address, and it can run a web search. Either can carry information from the conversation in the text of the request. An instruction that says "look up the following to help me" followed by content from the user's own mail produces an outbound request whose URL or query string is the leak. The SSRF guard prevents that request from reaching an internal address; it does nothing about a public one, because fetching public pages is the feature.
- **Write into memory.** The assistant can save a durable fact about its owner, and an injected instruction can put something there that persists into later conversations.
- **Add things.** Adding a to-do or a record is treated as easily reversible and is not gated the way a deletion is.

The worst realistic outcome, stated plainly: **content the assistant had access to can be sent to an attacker's server through a URL or a search query, without anybody approving anything**. That is the honest ceiling, and it should be the thing your review weighs, because it is the one Archie's architecture does not close.

## What actually reduces it

- **The no-tools rule while reading mail** (section 9) removes the channel entirely for the unattended mailbox lane, which is the highest-volume source of untrusted text most people have. This is a real mitigation and it is why that design exists.
- **Web access is not on by default.** Only a specialist can be given web access, and only if the user sets one up. An agent with no web-capable specialist and no page-fetch skill has no outbound leak channel of this kind.
- **Fewer connected accounts is less to leak.** The assistant can only leak what it can read.

### IF THIS IS THE FINDING YOUR REVIEW STOPS ON

That is a reasonable place to stop. The configuration that reduces it furthest is an agent with no web-fetch skill and no web-enabled specialist, which removes the outbound channel while leaving mail, calendar, and chat working. If your policy cannot accept a residual leak path from untrusted text, the answer is that Archie cannot satisfy it today, and no agent product we know of can. We would rather write that than have you find it yourself in month two.

# Add-ons, and what approving Archie approves

Archie has a marketplace of roughly 149 add-ons. Approving Archie means approving that supply chain, so here is what it actually is.

**An add-on is data, not code.** A skill is a text file plus some settings, a personality is written instructions, a routine is a schedule. There is no mechanism by which an add-on runs as a program, so an add-on cannot itself read your disk or open a socket. That is a stronger claim than sandboxing and it is true at the operating system layer.

WHERE THAT CLAIM STOPS BEING REASSURING

It is true at the operating system layer and misleading if it is left there. **An add-on is instructions to an agent that holds tools**, and text that steers a tool-holding agent is capability by another name. A malicious add-on has exactly the same leak channel as any hostile text the assistant reads, which is section 11. "It cannot run code" and "it cannot cause harm" are different sentences, and only the first is true.

Two further facts a reviewer should weigh:

- **Review is two people reading submissions by hand.** Every add-on in the catalog has been read line by line by one of the two people who run Otian AI. That is a real control at this size and it is not one that scales; the vendor says so on its own trust page rather than waiting to be asked.
- **Add-on permissions are blunter than they should be.** An add-on declaring "calendar" gets reading, creating, and deleting together, rather than being asked about each. The cap that saves this is the account grant itself: if the user connected their calendar read-only, no add-on gets create or delete no matter what it declared. So the meaningful control is the OAuth consent screen in section 9, not the add-on's declaration.

The practical consequence for a review: **the permission decision that matters is made once, at the consent screen, not per add-on.** If you grant read-only mail and read-only calendar, nothing installed later can exceed that.

# Telemetry, and the remote switches

Archie collects two things. This is the whole list.

| WHAT         | CONTENTS                                                                                                                                                   | HOW OFTEN                     |
|--------------|------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------|
| Heartbeat    | App version, platform, edition, and the time it was last seen. One record per account per edition.                                                         | At most once every six hours. |
| Error report | The tail of a crash, and whether the app is stuck restarting. Queued to disk so a crash can report itself on the next launch rather than during the crash. | Only when the app crashes.    |

There is no feature usage tracking, no analytics of any kind, no counts of messages or actions, and no message content. Every string in a report passes the same redactor that guards the logs. A desktop app is the one place a bad release cannot be pulled back, because the broken copy is on someone else's computer, so the vendor needs to know which versions are running and whether one is crashing.

**It can be turned off**, in Settings, under Crash reports. The same screen lists what a report contains. Turning it off stops the heartbeat, stops the uploads, and deletes anything already queued.

### The remote switches, and their authenticity story

The vendor can, without shipping a release, switch off a broken add-on, switch off a messaging channel, post a notice in the app, or set a minimum version below which agents refuse to start. This is a containment mechanism for a bad release. It cannot read anything, cannot reach a file, and cannot make the app do something it does not already do. Every default is permissive: a missing document, an absent field, and an unreachable server all behave identically, so a network that blocks the vendor entirely does not hand anyone a lever.

THE PART THAT IS NOT DEFENDED, AND SHOULD BE

The configuration document is fetched from Firestore over HTTPS and **is not signed**. Its authenticity rests on TLS, on Firestore's access rules, and on control of the vendor's Firebase project. Permissive defaults protect you against the document being *missing*; they do not protect you against it being *wrong*. The minimum-version field is the one that matters here, because its whole purpose is to stop agents starting, so a rules mistake or an account compromise at the vendor is a path to stopping every install from running agents.

It cannot exfiltrate, install, or execute anything: the worst case is denial of service plus a false notice in the app. Signing that document is on the list in section 17.

SECTION 14

## Permissions the operating system will ask about

Each of these produces a system prompt the user must accept. None is required for the app to run, and each corresponds to one feature.

| PROMPT        | PLATFORM           | WHAT IT IS FOR, AND WHAT HAPPENS IF IT IS DENIED                                                                                                                                                     |
|---------------|--------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Microphone    | macOS and Windows  | Talking to the assistant instead of typing. Speech is turned into text on this computer and never uploaded for transcription. Denied: the microphone button does not work, and nothing else changes. |
| Local network | macOS 15 and later | Finding smart lights and plugs on the user's own wifi. Denied: the device scan finds nothing, and nothing else changes. On a corporate network, denying it is a reasonable default.                  |



| PROMPT           | PLATFORM   | WHAT IT IS FOR, AND WHAT HAPPENS IF IT IS DENIED                                                                                                                                                                                                                                                                                                                                       |
|------------------|------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Full Disk Access | macOS only | Required only for the optional feature that reads the Mac's Messages database to propose replies to texts. The user must grant it by hand in System Settings; no application can grant it to itself. Denied, or simply not granted: the text feature does not run, and nothing else changes. This is a broad permission and it is worth withholding unless that one feature is wanted. |

Archie deliberately does not request Contacts access.

#### THE MACOS ENTITLEMENTS, WHICH ARE NOT A FORMALITY

The macOS build is signed with two hardened-runtime relaxations, `com.apple.security.cs.disable-library-validation` and `com.apple.security.cs.allow-dyld-environment-variables`. They exist because one optional bundled media tool is packaged with PyInstaller and unpacks its own Python libraries at launch, which library validation would otherwise reject. That is the standard entitlement set for shipping that kind of notarized binary, and saying so does not make it harmless.

What it means concretely: **library validation is off for the application process, and that process holds Keychain access**. Anything already able to write a dylib where the process will load it, or to set `DYLD_` variables in its environment, has an injection surface into a process holding your credentials. This does not help a remote attacker and it does not weaken anything against the network. It weakens the local case, which section 2 already lists as undefended, and it makes that case somewhat easier than it would otherwise be. The fix is to move the PyInstaller tool into a separate helper with its own unentitled signature, and it is on the list in section 17.

## SECTION 15

# Signing, distribution, and updates

| PLATFORM | SIGNED             | DETAILS                                                                                                                                                                                                                                                                                                                               |
|----------|--------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| macOS    | Yes, and notarized | Signed with an Apple Developer ID certificate issued to Jett Nguyen, team ID <code>5SLXY6F2CD</code> , then submitted to Apple for notarization and stapled. Gatekeeper opens it without a warning. The signature can be checked with <code>codesign -dv --verbose=4</code> and <code>spctl -a -vv</code> against the downloaded app. |
| Windows  | No                 | The installer is not code-signed, so SmartScreen shows a warning on first launch. See the note below.                                                                                                                                                                                                                                 |

### STATED PLAINLY, BECAUSE YOU WOULD FIND IT ANYWAY

The Windows installer ships unsigned. This is a gap and it is known. It is not a cost decision: the affordable route to a code-signing certificate requires three years of company history, and Otian AI LLC is not old enough yet. The practical effect is that a Windows user meets a SmartScreen warning and has to click through it, which is exactly the kind of thing an IT department is right to treat as a red flag when it appears unexplained.

What can be verified instead: the download comes from a known GitHub release under the publisher's own account, the file hash can be compared with the published release, and the macOS build of the same application is fully signed and notarized by Apple. If an unsigned installer is not acceptable under policy, the macOS build is the one to approve.

## Updates, and what the signature proves

Application updates are checked at launch against a manifest served from `otianai.com`, downloaded from GitHub, and verified against a minisign public key compiled into the app before they are applied. An update that fails that check is not installed. The update is offered to the user rather than applied silently, and requires no administrator rights.

Three things that signature does not prove, and a reviewer should have all three:

- **It does not prove the build is current.** The signature says the package came from us. An older, validly signed, vulnerable build carries a valid signature forever. Protection against a downgrade rests on the manifest, and the manifest is protected by TLS and by control of `otianai.com`, not by a signature of its own. During the beta that manifest also sits at an unlisted URL rather than behind an access check, which is obscurity and not access control.
- **The signing key is held by a two-person company.** It lives in the release environment of the person who cuts builds. There is no hardware security module, no threshold scheme, and no published rotation or revocation procedure. If that key were compromised, the recovery path today is shipping a new key in a new build, which is exactly the circular problem it sounds like.
- **The updater is the one component that cannot be fixed remotely.** A broken updater means contacting users individually, which is why the release process treats it as the highest-risk change.

### SECTION 16

## What the AI provider does with the data

---

Archie does not choose this. The user pastes an API key for a provider they have an account with, and that provider's terms govern what happens to the text.

Because that difference lands on whoever pasted the key, each provider's own wording is kept on file, quoted rather than paraphrased, with the page it came from and the date it was read. Twice-yearly re-reading is part of the arrangement, because a stale quote on a page whose whole worth is that it can be checked is worse than no quote. This is the current summary, read on August 16, 2026:

| PROVIDER           | TRAINS ON WHAT YOU SEND?                                                                                                                                                                                | THEIR OWN PAGE                                                                                                                                  |
|--------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------|
| Claude (Anthropic) | No                                                                                                                                                                                                      | <a href="https://anthropic.com/legal/commercial-terms">anthropic.com/legal/commercial-terms</a>                                                 |
| ChatGPT (OpenAI)   | No, unless you opt in. Abuse-monitoring logs are kept up to 30 days.                                                                                                                                    | <a href="https://developers.openai.com/api/docs/guides/your-data">developers.openai.com/api/docs/guides/your-data</a>                           |
| Groq               | No                                                                                                                                                                                                      | <a href="https://console.groq.com/docs/legal/services-agreement">console.groq.com/docs/legal/services-agreement</a>                             |
| Gemini (Google)    | Free key: yes. Paid key: no. An AI Studio key can be either, and Archie has never asked which.                                                                                                          | <a href="https://ai.google.dev/gemini-api/terms">ai.google.dev/gemini-api/terms</a>                                                             |
| Mistral            | Free mode: yes by default. Pay as you go: no.                                                                                                                                                           | <a href="https://legal.mistral.ai/terms/commercial-terms-of-service">legal.mistral.ai/terms/commercial-terms-of-service</a>                     |
| DeepSeek           | Yes, and the data is stored in China, if you use DeepSeek's own hosted service. Their models are open weight, so a copy running on hardware you control sends them nothing.                             | <a href="https://cdn.deepseek.com/policies/en-US/deepseek-privacy-policy.html">cdn.deepseek.com/policies/en-US/deepseek-privacy-policy.html</a> |
| Grok (xAI)         | Unverified. Their primary terms return an error to every automated reader, so there is no quote to stand behind. What is corroborated is that the consumer Grok app trains on conversations by default. | <a href="https://x.ai/legal">x.ai/legal</a> (unreadable as of the date above)                                                                   |

## How an API key compares with an enterprise plan

This is the question most reviewers actually want answered, and the honest answer has two halves.

**On what happens to the text, an API key gives you what an enterprise plan gives you.** Anthropic, OpenAI, and Groq all state that content sent through their API is not used to train models, and that is the default for an ordinary paid key rather than a term somebody has to negotiate. It is the consumer chat subscription that is the weaker product here.

**On what you can prove and enforce, it does not.** An enterprise or business agreement is a contract, and it carries things a self-service API key does not:

| WHAT ENTERPRISE ADDS               | WHY AN API KEY ALONE DOES NOT HAVE IT                                                                                                                                             |
|------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| A signed data processing agreement | Standard API terms are accepted by clicking, not signed. If your policy requires a countersigned DPA naming your organization, that comes with the commercial agreement.          |
| Zero data retention                | OpenAI keeps abuse-monitoring logs for up to 30 days on ordinary API use. Removing that is a separate arrangement you apply for, not a default.                                   |
| Compliance artifacts               | SOC 2 reports, a HIPAA business associate agreement, and data residency commitments are provided under a commercial agreement, not attached to a key you generated on a web form. |

| WHAT ENTERPRISE ADDS   | WHY AN API KEY ALONE DOES NOT HAVE IT                                                                                                                                                           |
|------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Administrative control | Single sign-on, seat management, spending limits, and the ability to revoke one person's access centrally belong to the organization account. A personal key belongs to the person who made it. |

THE PRACTICAL ANSWER FOR AN ORGANIZATION

If your organization already has an enterprise or team agreement with Anthropic, OpenAI, or another provider, issue an API key from that account and have the user paste that key. Archie treats it like any other key, the traffic goes straight from the computer to your provider, and every guarantee in your existing contract applies to it without Archie being party to any of it. That is the strongest configuration available and it needs nothing special on Archie's side.

Running the model on your own hardware

Archie also offers a provider option called **Another provider**, which takes any endpoint speaking the widely used OpenAI-compatible format along with the model name to ask for. That covers a model your organization hosts itself. Under that arrangement the text never reaches any AI vendor: it goes from this computer to a machine you operate. For a team that cannot send its work to a third party at all, this is the route.

SECTION 17

Open points, stated rather than buried

Archie is at version 0.1.7 and is in beta. A document that lists only the good parts is not useful to somebody deciding whether to trust it, so here is the rest, including the things that have their own callouts earlier.

| POINT                                           | WHAT IT MEANS FOR A REVIEWER                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
|-------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Prompt injection can leak                       | The largest one, and it has its own section (11). An injected instruction cannot change your accounts silently, and it can carry conversation content out through a fetched URL or a search query.                                                                                                                                                                                                                                                                                                        |
| The trial mail gate is a name-prefix denylist   | Section 8. It covers every mail tool that exists, and it depends on a naming convention rather than a declared capability. Moving it to a capability tag is the fix and it has not been done.                                                                                                                                                                                                                                                                                                             |
| The Windows installer is unsigned               | Section 15. Users meet a SmartScreen warning. The macOS build is signed and notarized.                                                                                                                                                                                                                                                                                                                                                                                                                    |
| Windows credential storage is weaker than macOS | Keychain items carry per-application access control lists. Windows Credential Manager gives user-scoped DPAPI protection with no cryptographic binding to a signed binary. Combined with an unsigned installer in a user-writable location, any process running as that user can replace the Archie executable and inherit access to its stored credentials. On Windows, treat credential protection as "protected from other users on this machine", not "protected from other software this user runs". |

| POINT                                                                | WHAT IT MEANS FOR A REVIEWER                                                                                                                                                                                                              |
|----------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>The macOS entitlements widen the local surface</b>                | Section 14. Library validation is disabled on a process holding Keychain access, to support one optional bundled tool. Splitting that tool into an unentitled helper is the fix and has not been done.                                    |
| <b>The remote configuration document is unsigned</b>                 | Section 13. Worst case is denial of service and a false in-app notice, not code execution.                                                                                                                                                |
| <b>No downgrade protection, and single-person key custody</b>        | Section 15. The signature proves origin, not recency, and the manifest that establishes recency is protected by TLS and host control alone.                                                                                               |
| <b>The audit log is not attacker-proof</b>                           | Section 10. Tamper-evident against corruption and casual editing; the anchor is reachable by the same session.                                                                                                                            |
| <b>Chat platforms are a real egress path</b>                         | Section 7. Optional, and the strongest answer is not connecting one, or using a Matrix homeserver you run.                                                                                                                                |
| <b>The local database is not separately encrypted</b>                | Conversation history sits in a SQLite file readable by anyone with access to that user's profile. No credential is in it. Whole-disk encryption is the control.                                                                           |
| <b>Mail scopes are real scopes</b>                                   | Section 9. The propose-only behavior is enforced by the application; the token permits what the user granted.                                                                                                                             |
| <b>Add-on permissions are coarse</b>                                 | Section 12. Declaring "calendar" takes read, create, and delete together; the OAuth grant is the real cap.                                                                                                                                |
| <b>Some add-ons run helper binaries</b>                              | Optional media and speech tools run as child processes with an explicit argument list, a deadline, and a scratch directory deleted afterward. They are hash-verified before they can run. They are not sandboxed by the operating system. |
| <b>It is beta software and has never been audited by anyone else</b> | Version 0.1.7. Every feature described here works and has been exercised. There has been no external penetration test and no third-party security audit. Everything in this document is the vendor checking its own work.                 |

## SECTION 18

# The paperwork: what exists and what does not

A review usually asks for a set of documents around hour two. Here is which of them exist, so nobody spends a week discovering the answer.

| ARTIFACT                 | EXISTS     | WHERE, OR WHAT IS MISSING                                                                                                                                                                                                                                                                                                                                            |
|--------------------------|------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Subprocessor list</b> | <b>Yes</b> | Named in the privacy policy at <a href="https://otianai.com/privacy-policy">otianai.com/privacy-policy</a> : Google Firebase, Stripe, Resend, and Render. Render is the backend that handles sign-in and licensing and that relays AI requests while a user is on the free credits. Your AI provider is not a subprocessor of ours: you contract with them directly. |

| ARTIFACT                                    | EXISTS        | WHERE, OR WHAT IS MISSING                                                                                                                                                                                                                                                                                    |
|---------------------------------------------|---------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| GDPR and US state privacy rights            | Yes           | The privacy policy states the lawful bases relied on in the EEA and UK (contract, legitimate interests, consent), the access, correction, deletion, portability, and objection rights, and that no personal information is sold or shared for cross-context behavioral advertising.                          |
| Account deletion                            | Yes           | Self-service from the account page after a recent sign-in. It removes sign-in, account record, and purchase history. Crash reports, the version line, and the free-credit ledger with its abuse counters outlive it, and none of those contains anything the user wrote. They can be deleted on request.     |
| Retention schedule with numbers             | No            | The privacy policy commits to keeping records only as long as needed and deleting or anonymizing them after. There are no published retention periods for crash tails, invoices, or sealed phone messages. If your policy requires stated periods, ask and we will give you the current practice in writing. |
| Vulnerability disclosure policy             | No            | There is no published policy, no security advisory feed, and no dedicated security mailbox. Reports currently go to the contact address on this document. If you find something, that address reaches a person the same day.                                                                                 |
| Incident and breach notification commitment | No            | Nothing is published committing to a notification window. What limits the blast radius is the holdings list in section 7: a breach of Otian AI cannot produce anybody's conversations, mail, calendar, files, or keys, because none of those are held.                                                       |
| SOC 2, ISO 27001, HIPAA BAA                 | No            | None held, and none in progress. A two-person company at version 0.1.7. If your policy requires a certification from the software vendor, Archie does not meet it today, and no amount of architecture in this document changes that answer.                                                                 |
| Dependency inventory and CVE handling       | Not published | Dependencies are pinned in checked-in lock files for both the Rust and TypeScript sides, and four packages are deliberately held below latest with the blocker recorded for each. There is no published SBOM and no stated CVE response window. An SBOM can be produced on request from the lock files.      |
| External security audit                     | No            | Nobody has audited Archie but the people who wrote it.                                                                                                                                                                                                                                                       |

## REMOVING IT

Archie is a normal application and uninstalls like one. Deleting the app removes the program; deleting the folder named in section 4 removes its data; the credentials it created appear under `com.archie.app` in Keychain Access or Windows Credential Manager and can be deleted there. Access already granted to a mailbox or calendar should be revoked at the provider, in the Google account permissions page or the Microsoft 365 admin center, which is where OAuth grants actually live.

---

**Questions.** Anything in this document can be checked further, including specific network captures, the signing details on a given build, or a walkthrough of a feature before it is approved. If your review stopped on one row, say which one: we would rather answer that than send more pages. Write to Jack at [jack@otianai.com](mailto:jack@otianai.com).

Archie 0.1.7 · Otian AI LLC · [otianai.com](https://otianai.com) · Prepared August 25, 2026. Details describe this version and may change in later releases.